

# **CS 4530 & CS 5500** **Software Engineering**

## **Lecture 9.2: Strategies for Engineering Distributed Software**

**Jonathan Bell, John Boyland, Mitch Wand**  
**Khoury College of Computer Sciences**  
**© 2021, released under [CC BY-SA](#)**

# **Learning Objectives for this Lesson**

**By the end of this lesson, you should be able to...**

- Characterize the benefits of replication and partitioning in distributed systems
- Evaluate the tradeoffs between consistency and availability in distributed systems

# Why expand to distributed systems?

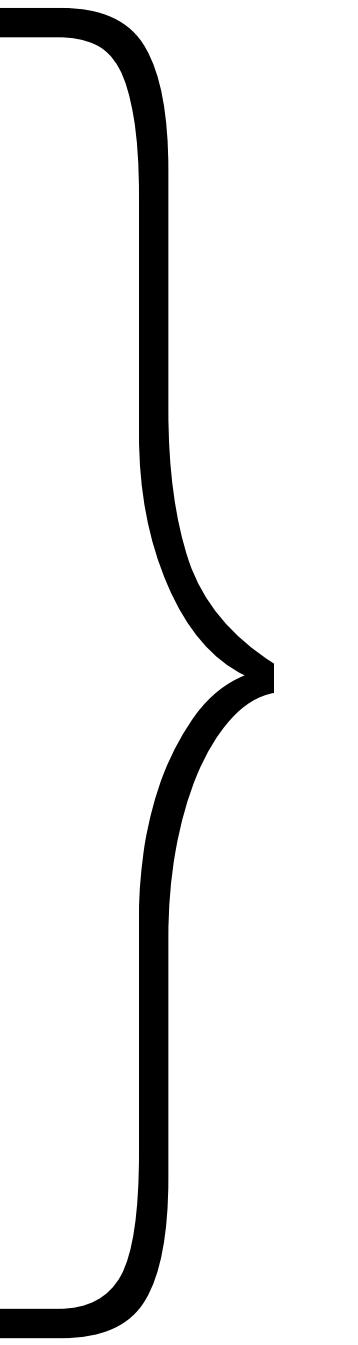
- Scalability
- Performance
- Latency
- Availability
- Fault Tolerance

“Distributed Systems for Fun and Profit”, Takada

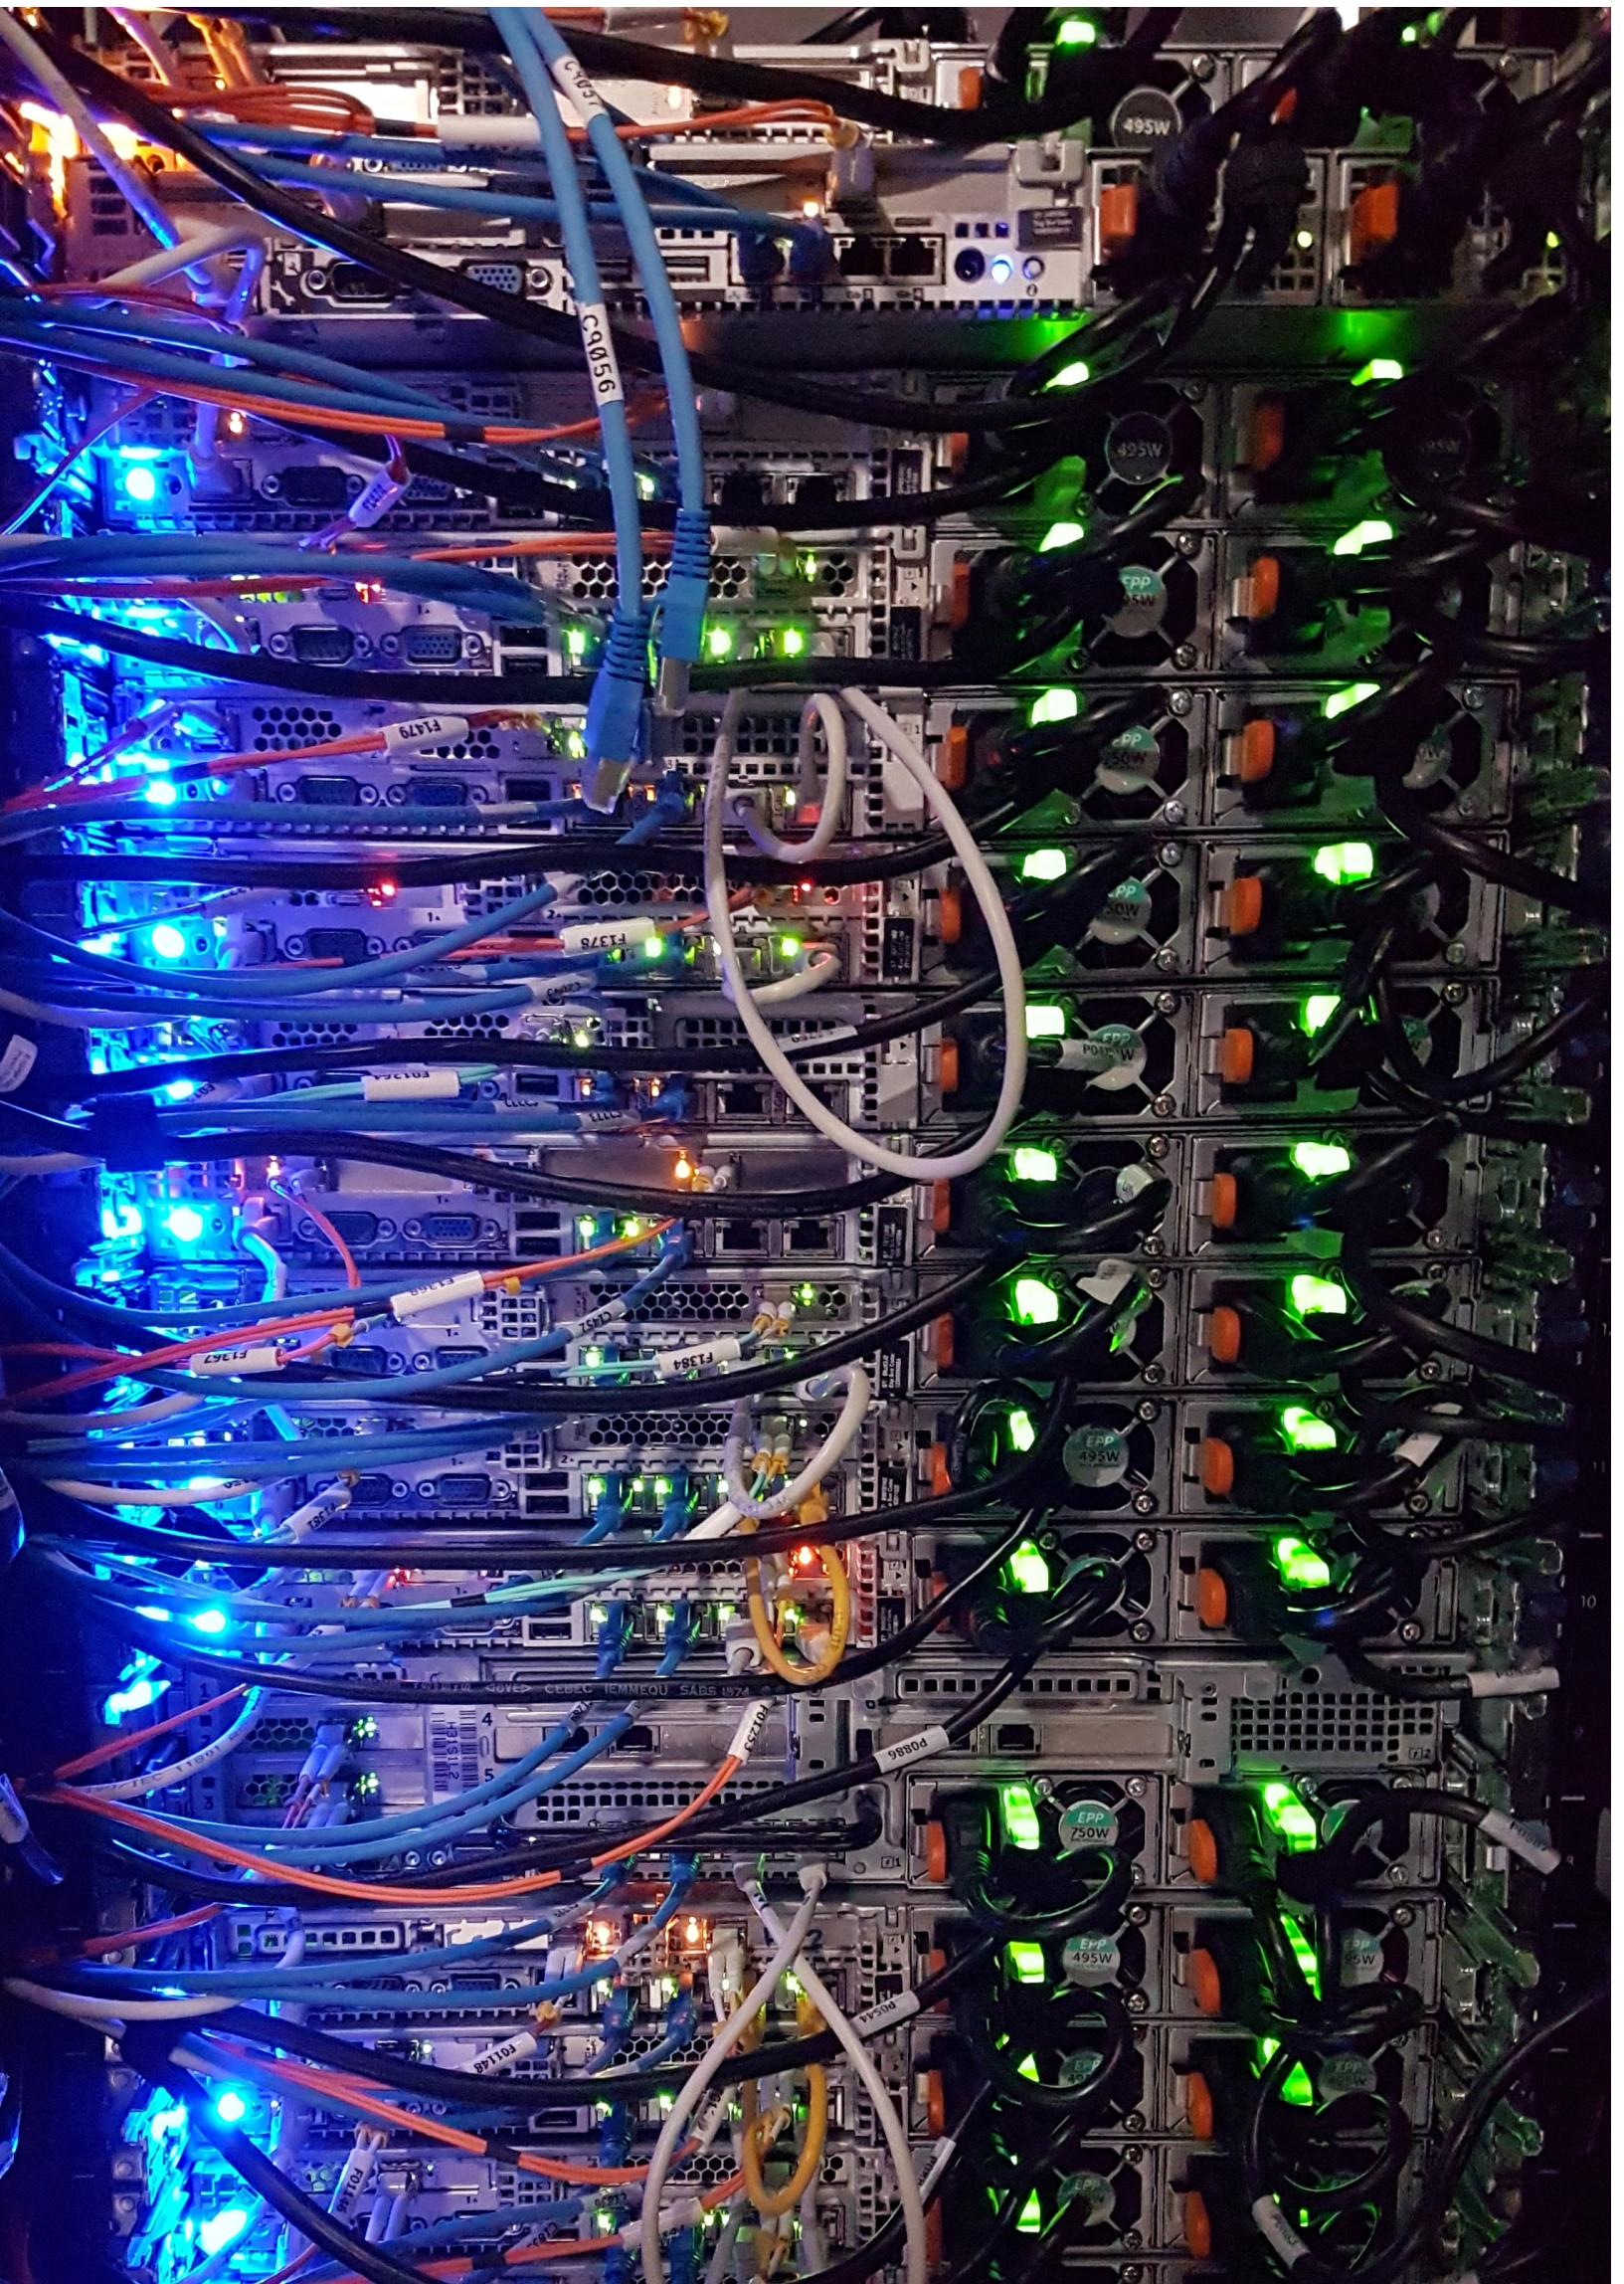
# 8 Fallacies of Distributed Computing

Sun Microsystems, early 1990's

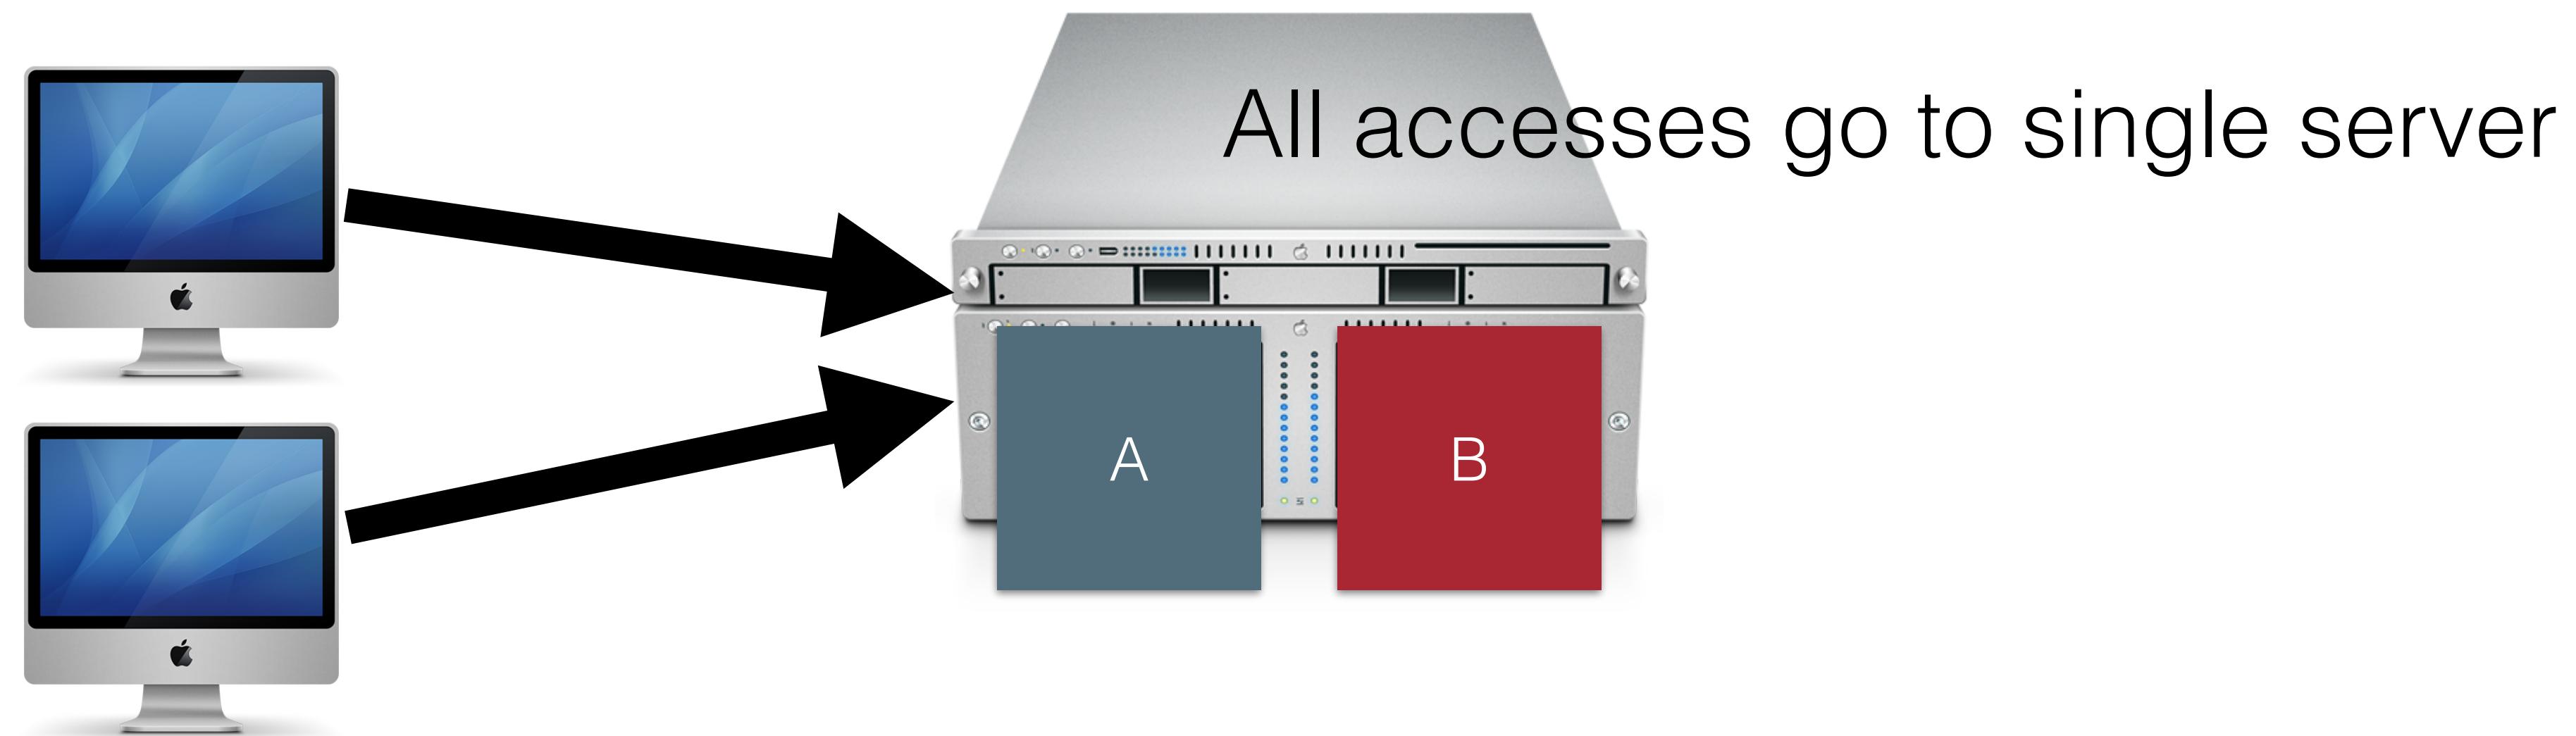
- 1. The network is reliable
- 2. Latency is zero
- 3. Bandwidth is infinite
- 4. The network is secure
- 5. Topology doesn't change
- 6. Transport cost is zero
- 7. The network is homogeneous
- 8. There is one administrator



Physical Properties of networks

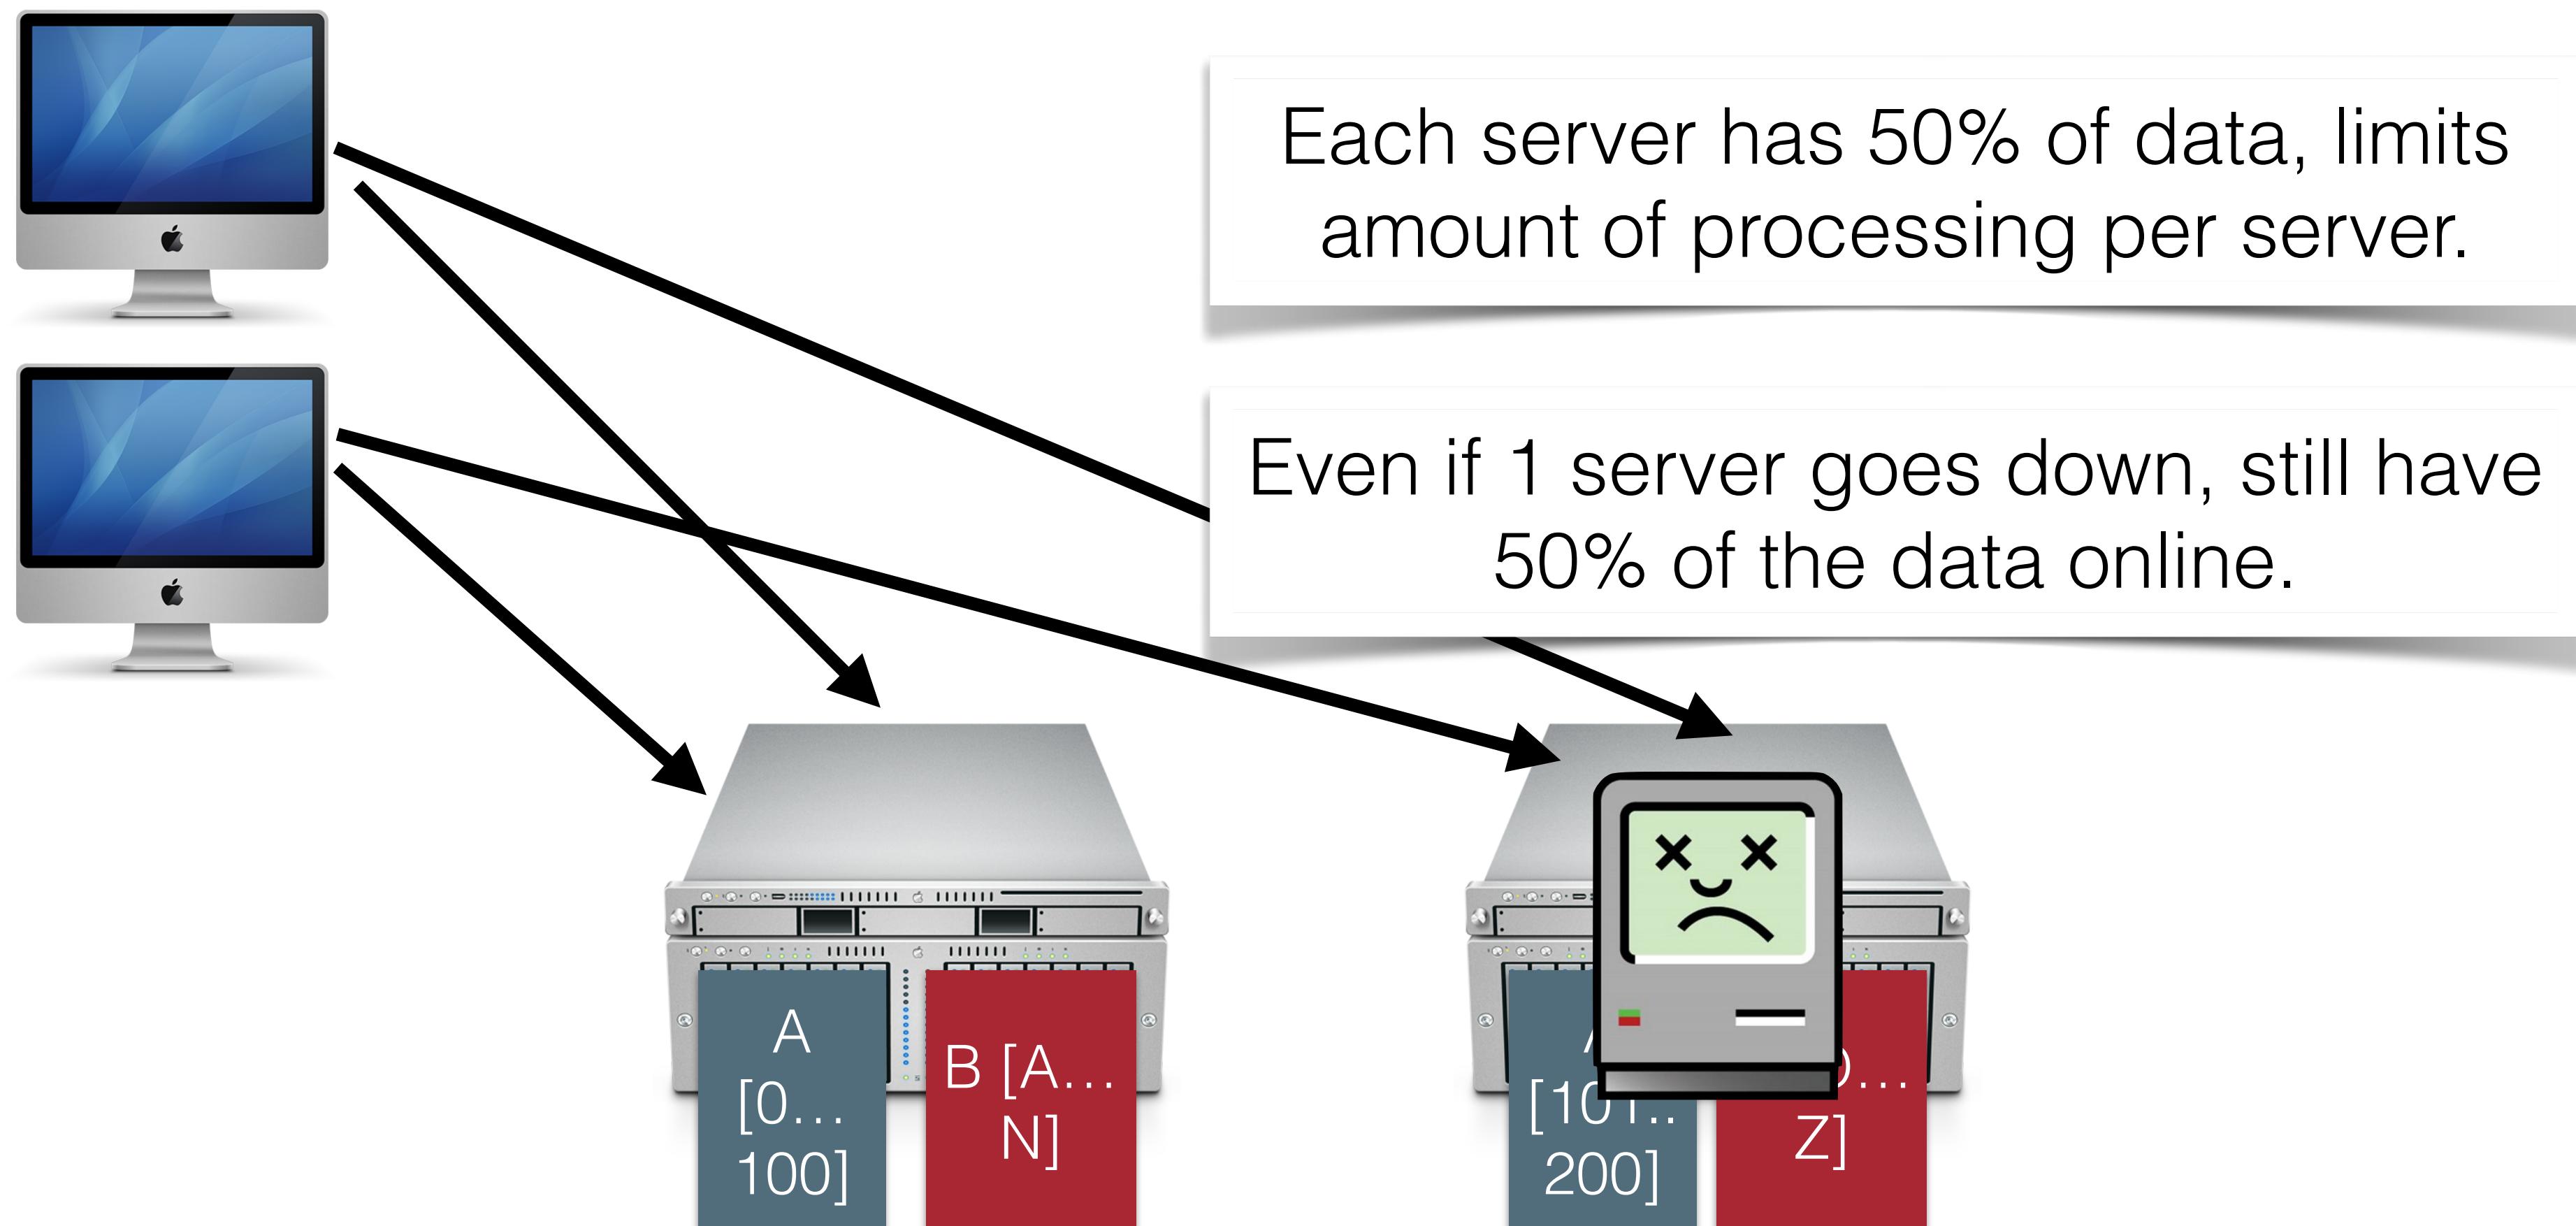


# Recurring Solution #1: Partitioning

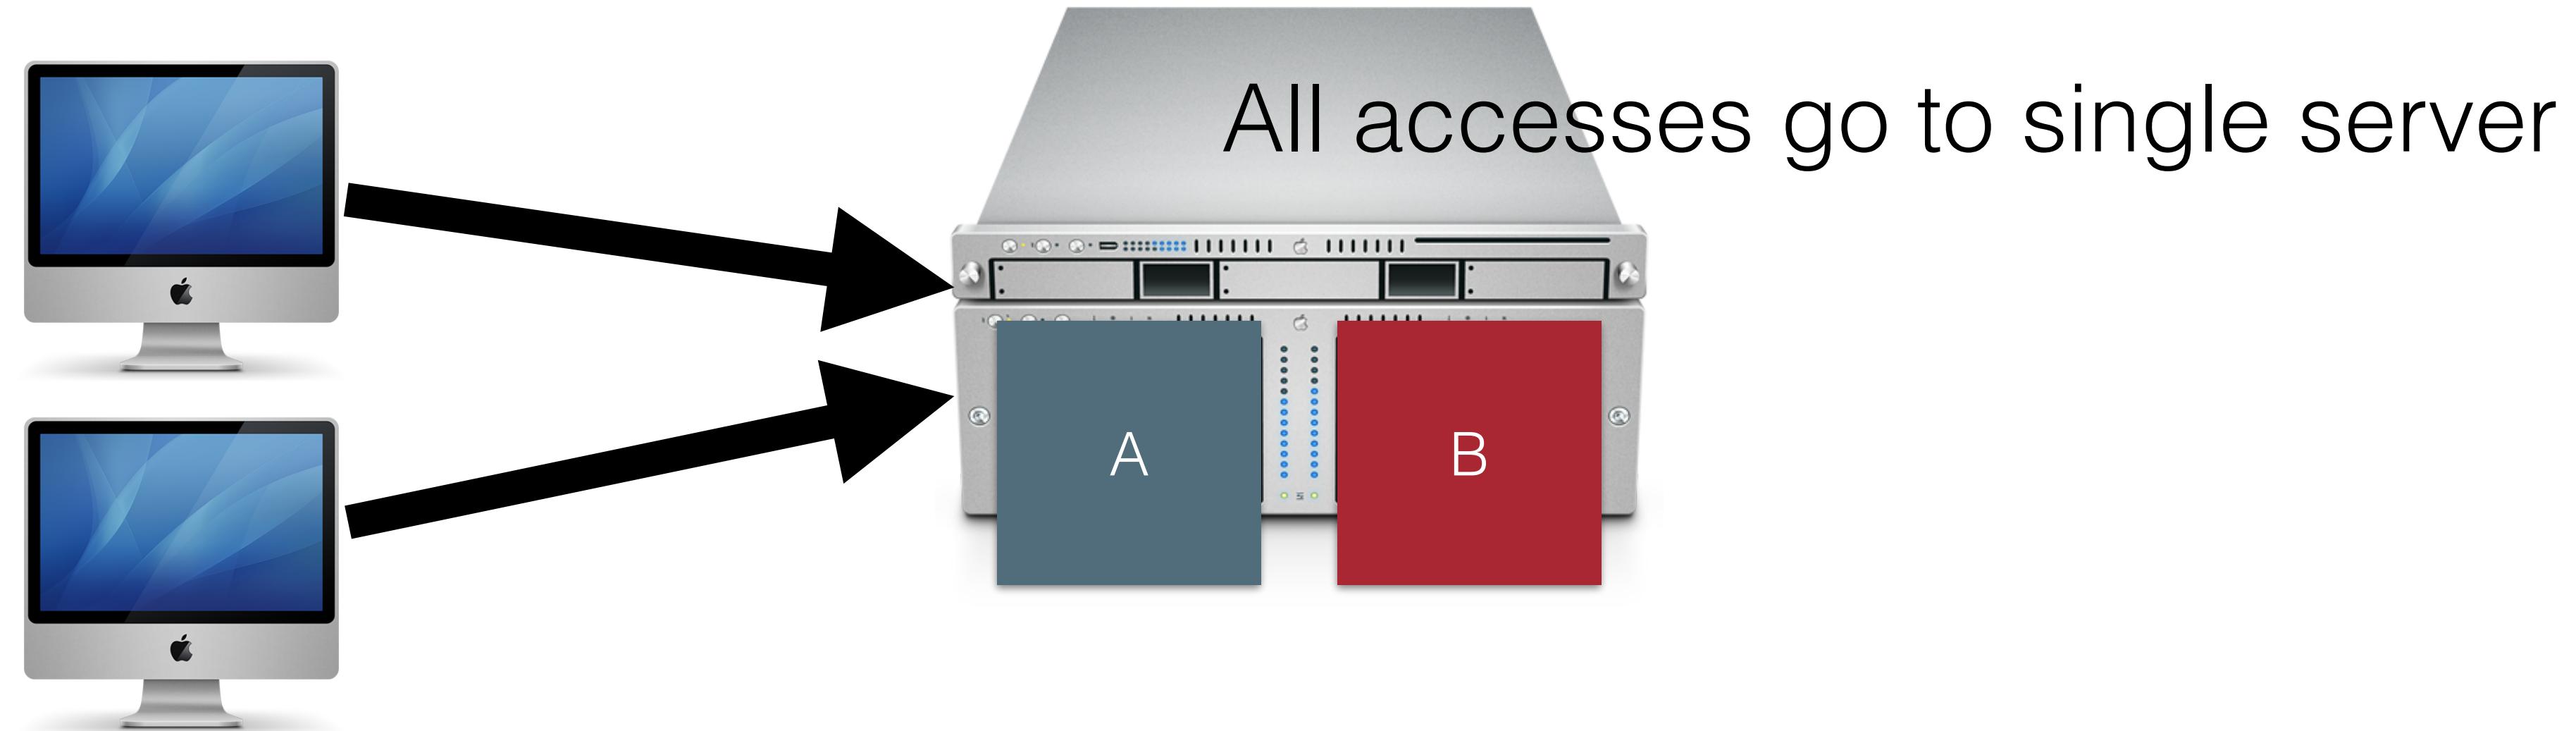


# Recurring Solution #1: Partitioning

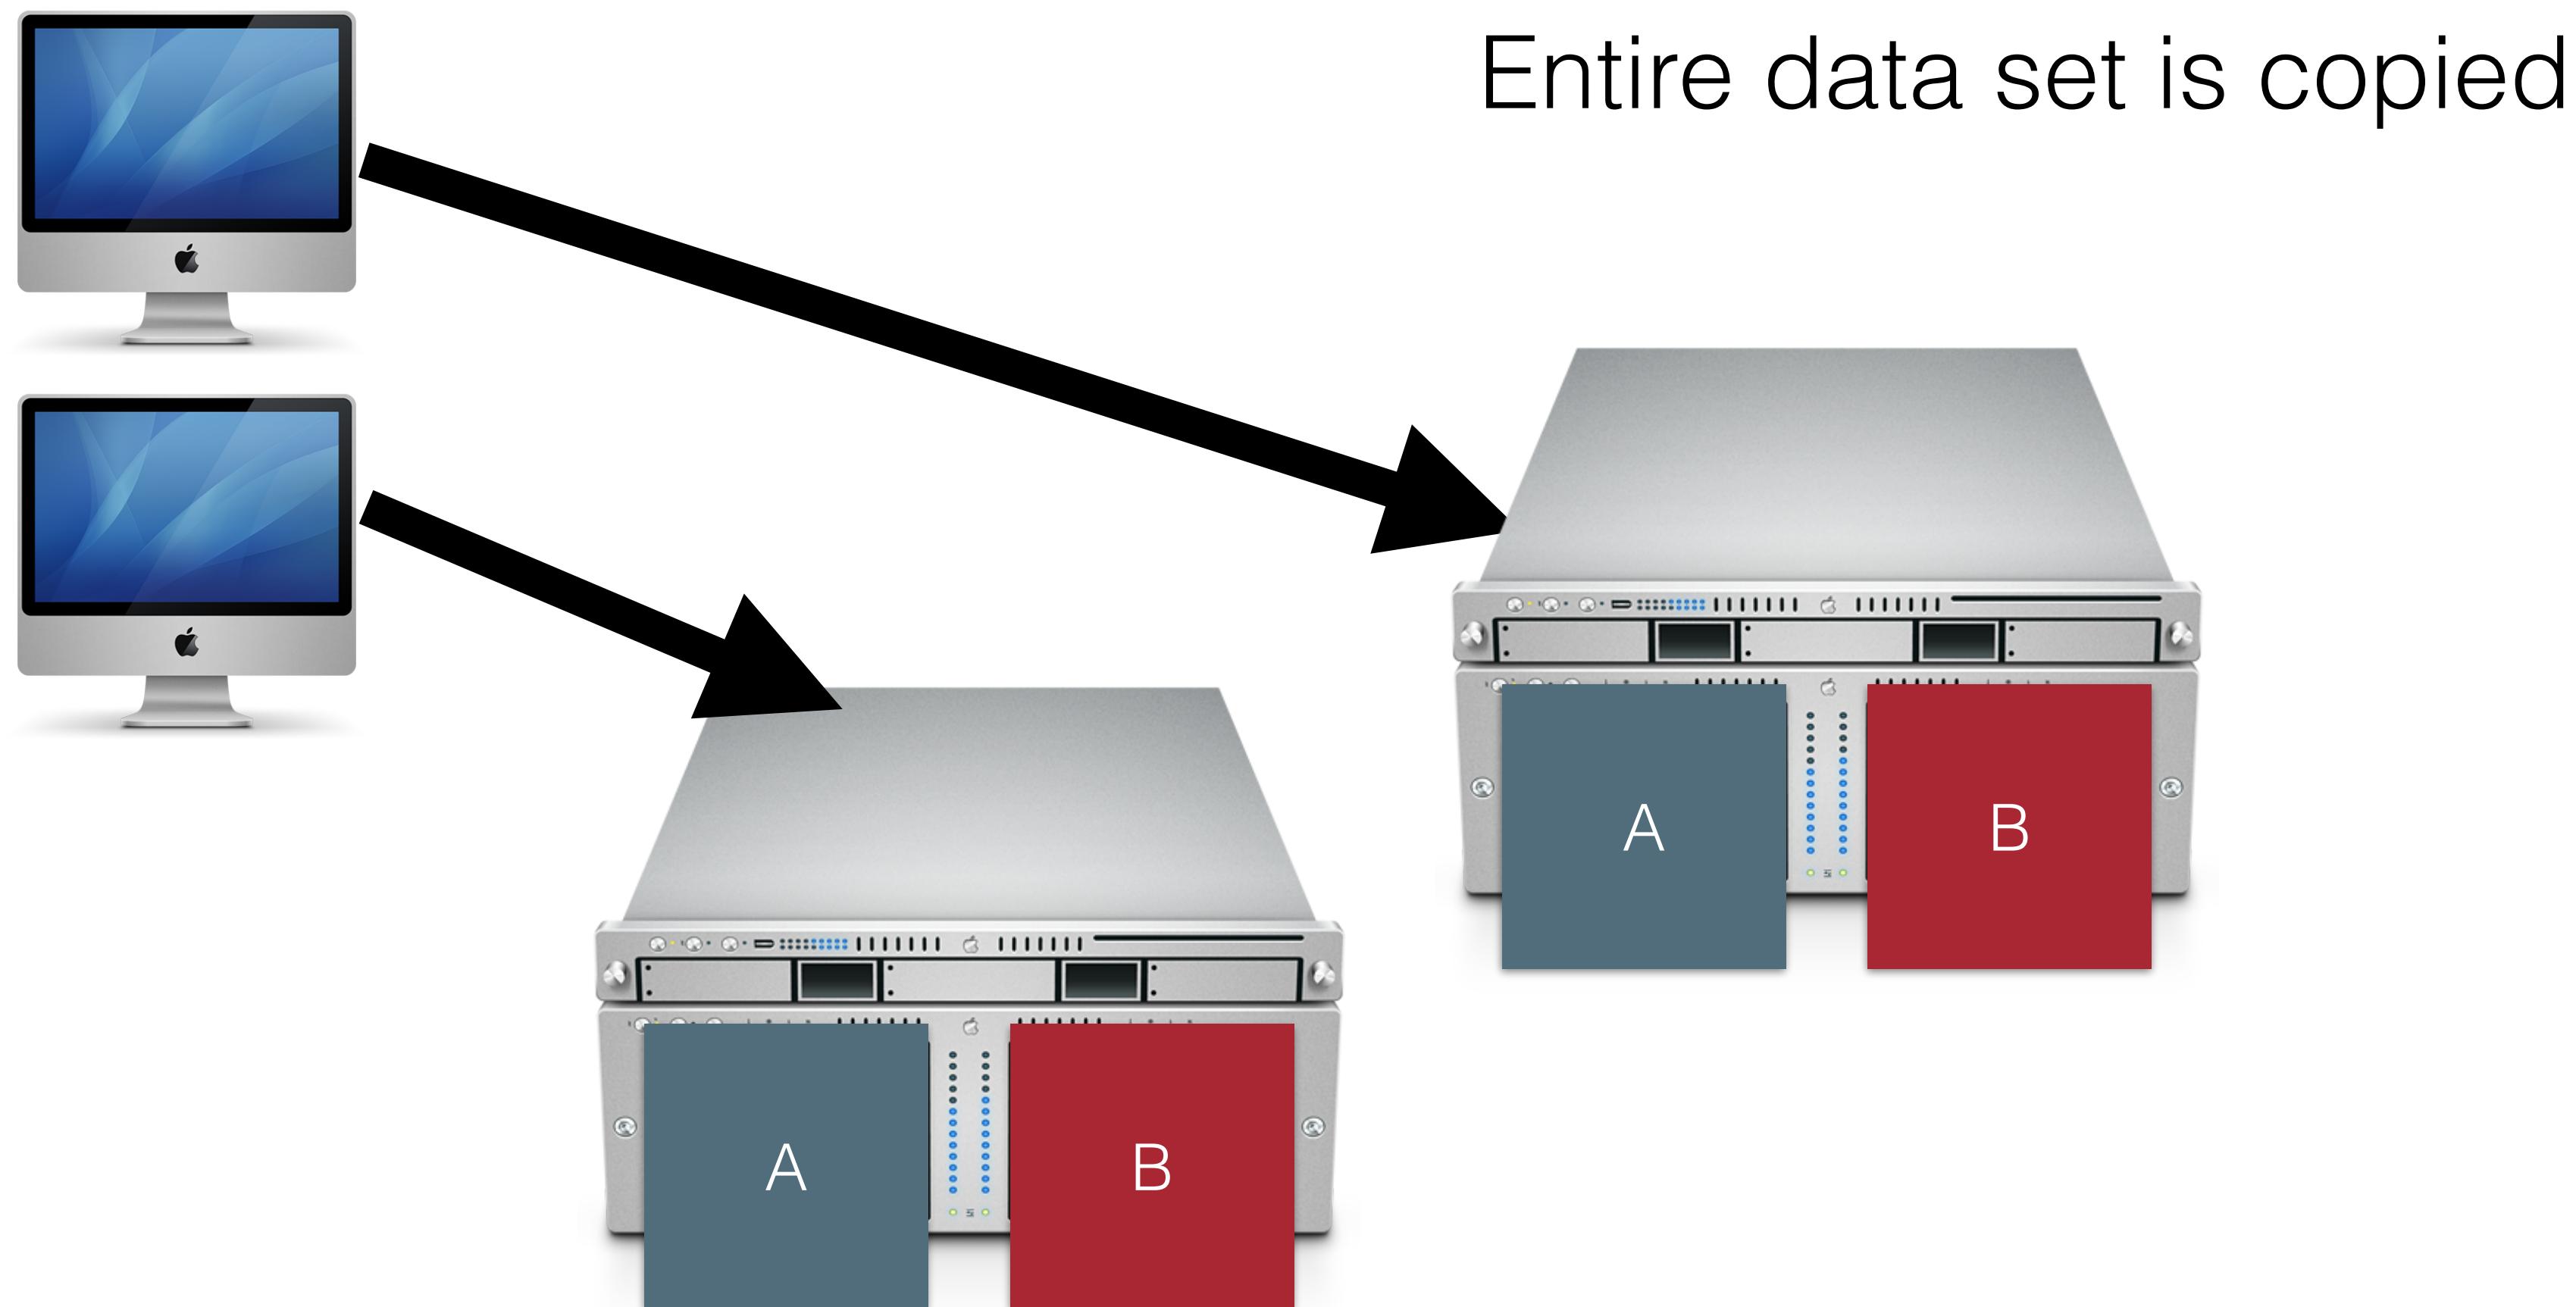
- Divide data up in some (hopefully logical) way
- Makes it easier to process data concurrently (cheaper reads)



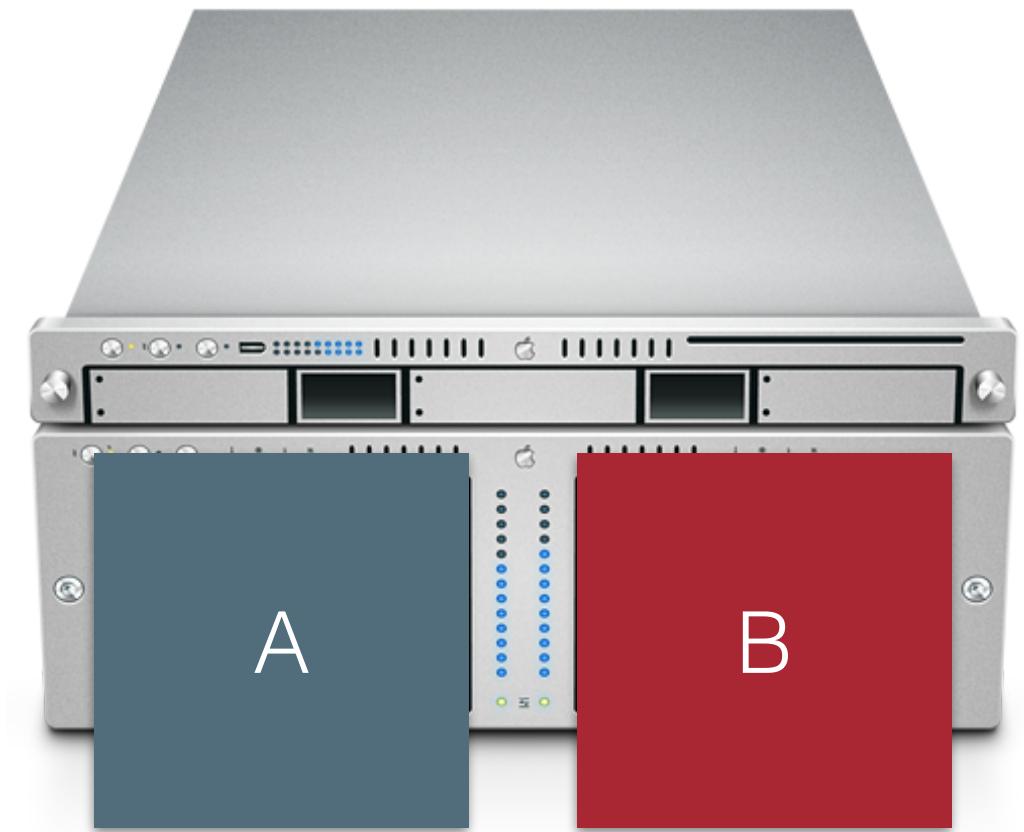
# Recurring Solution #2: Replication



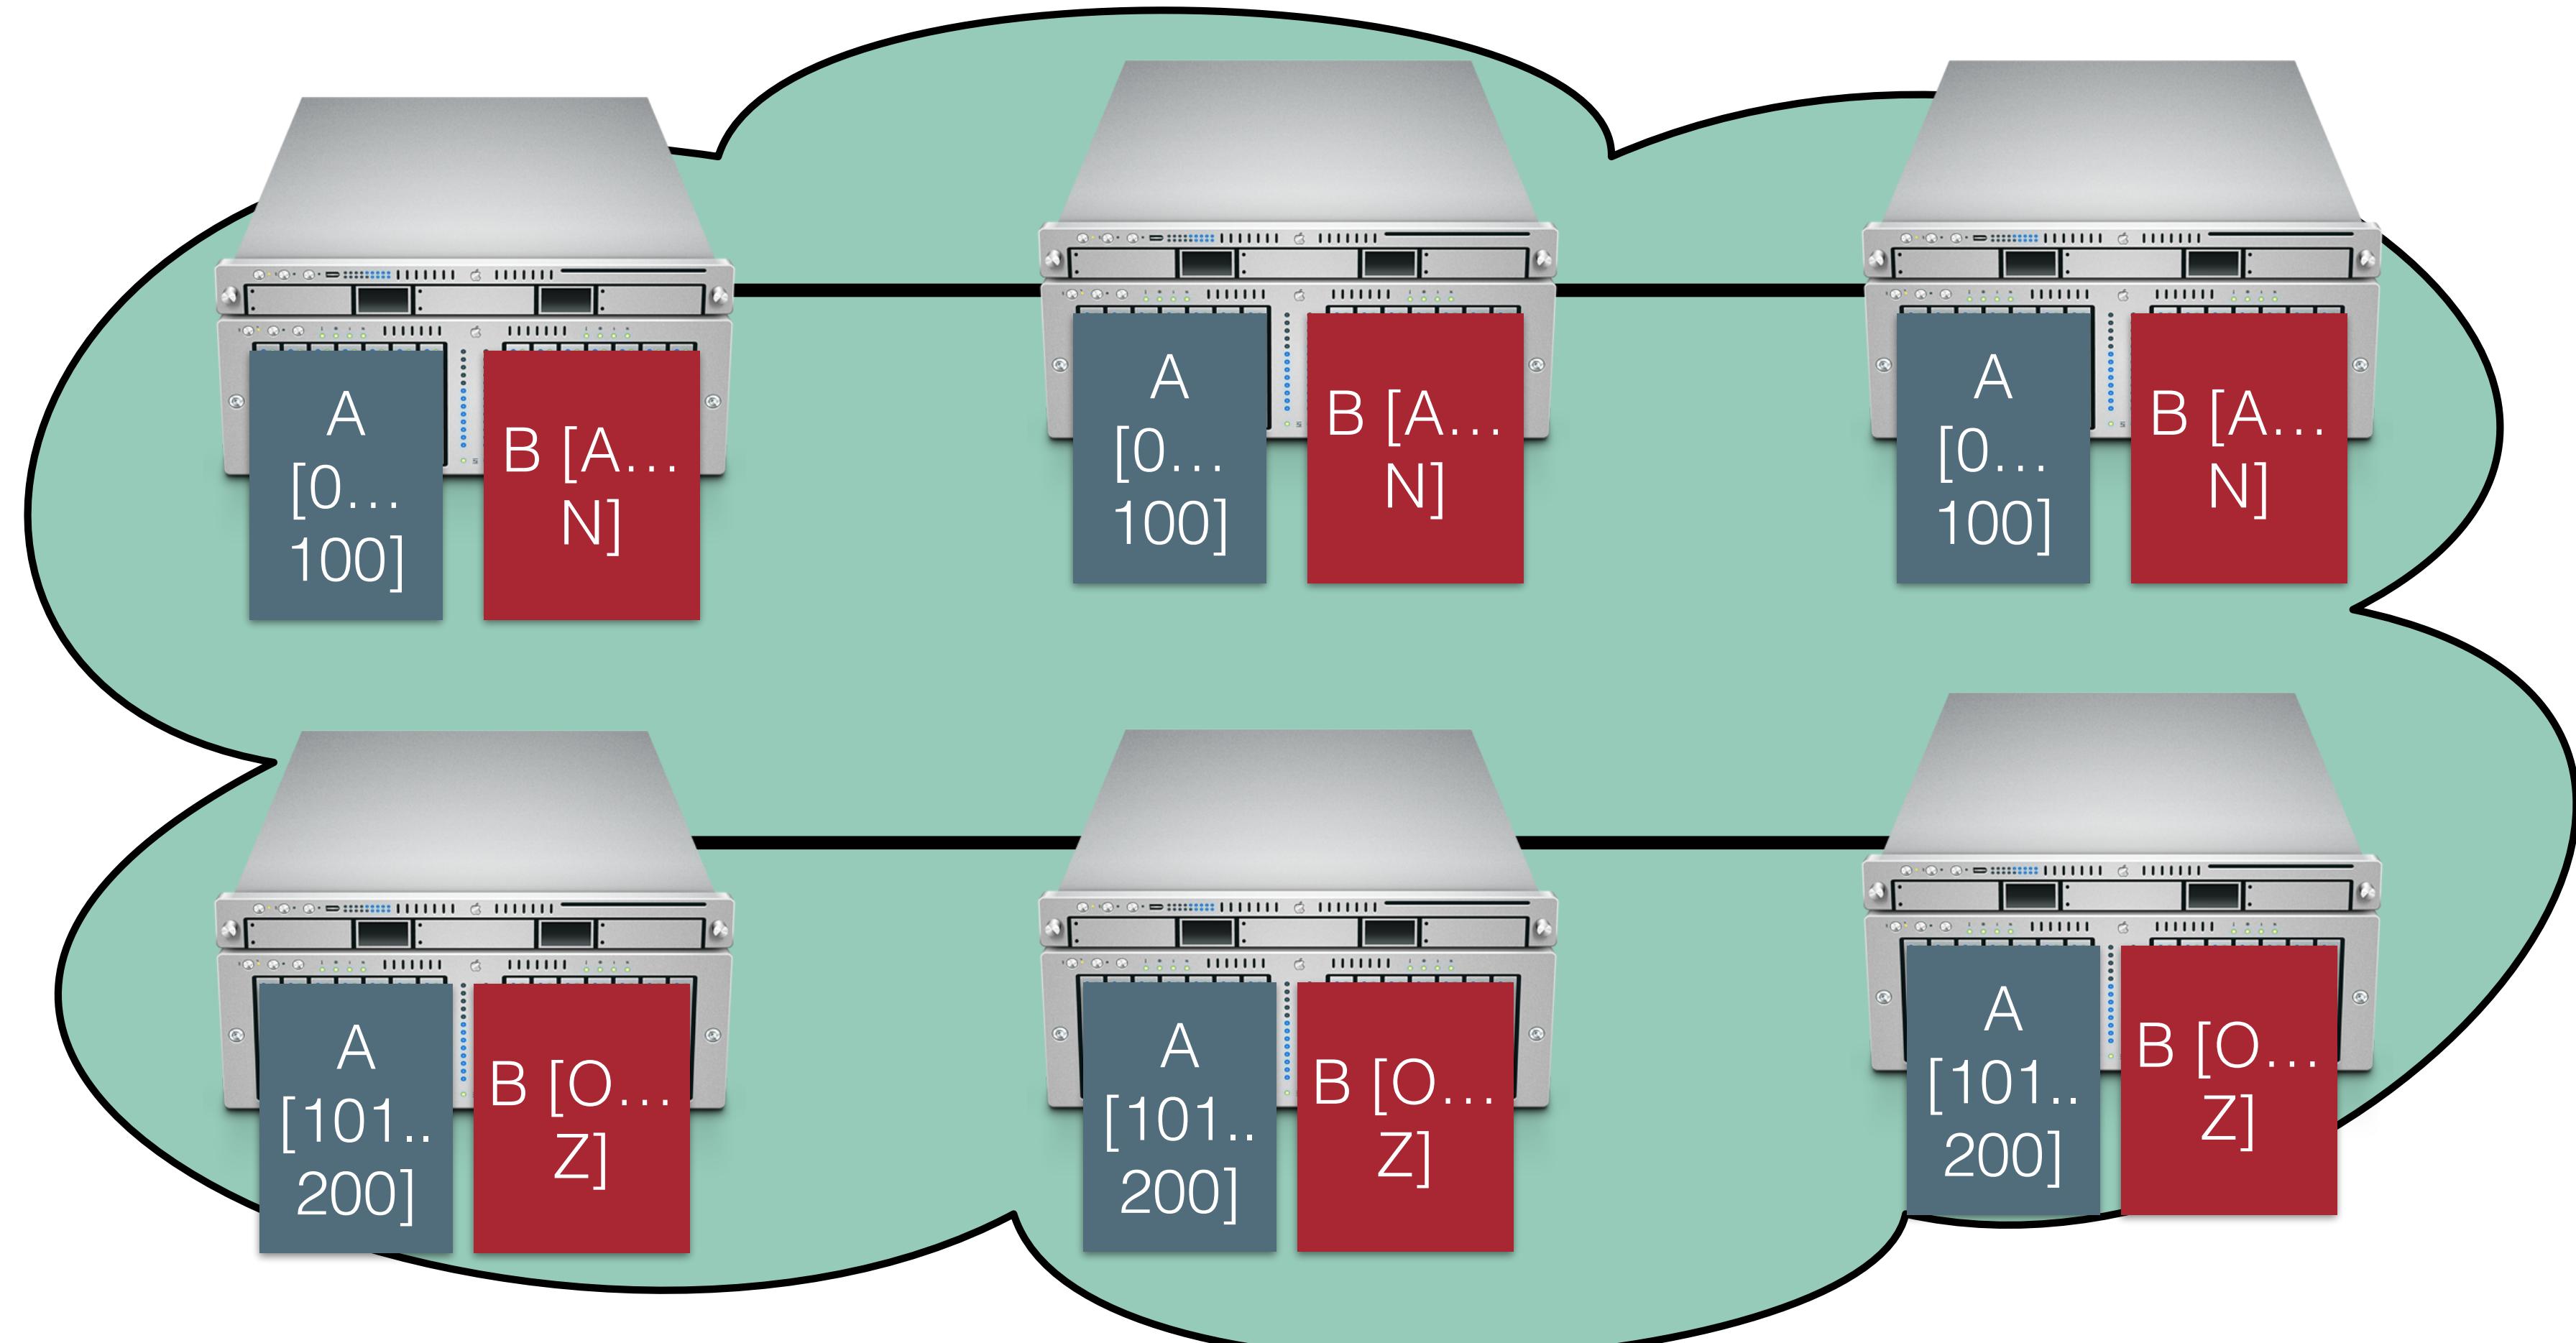
# Recurring Solution #2: Replication



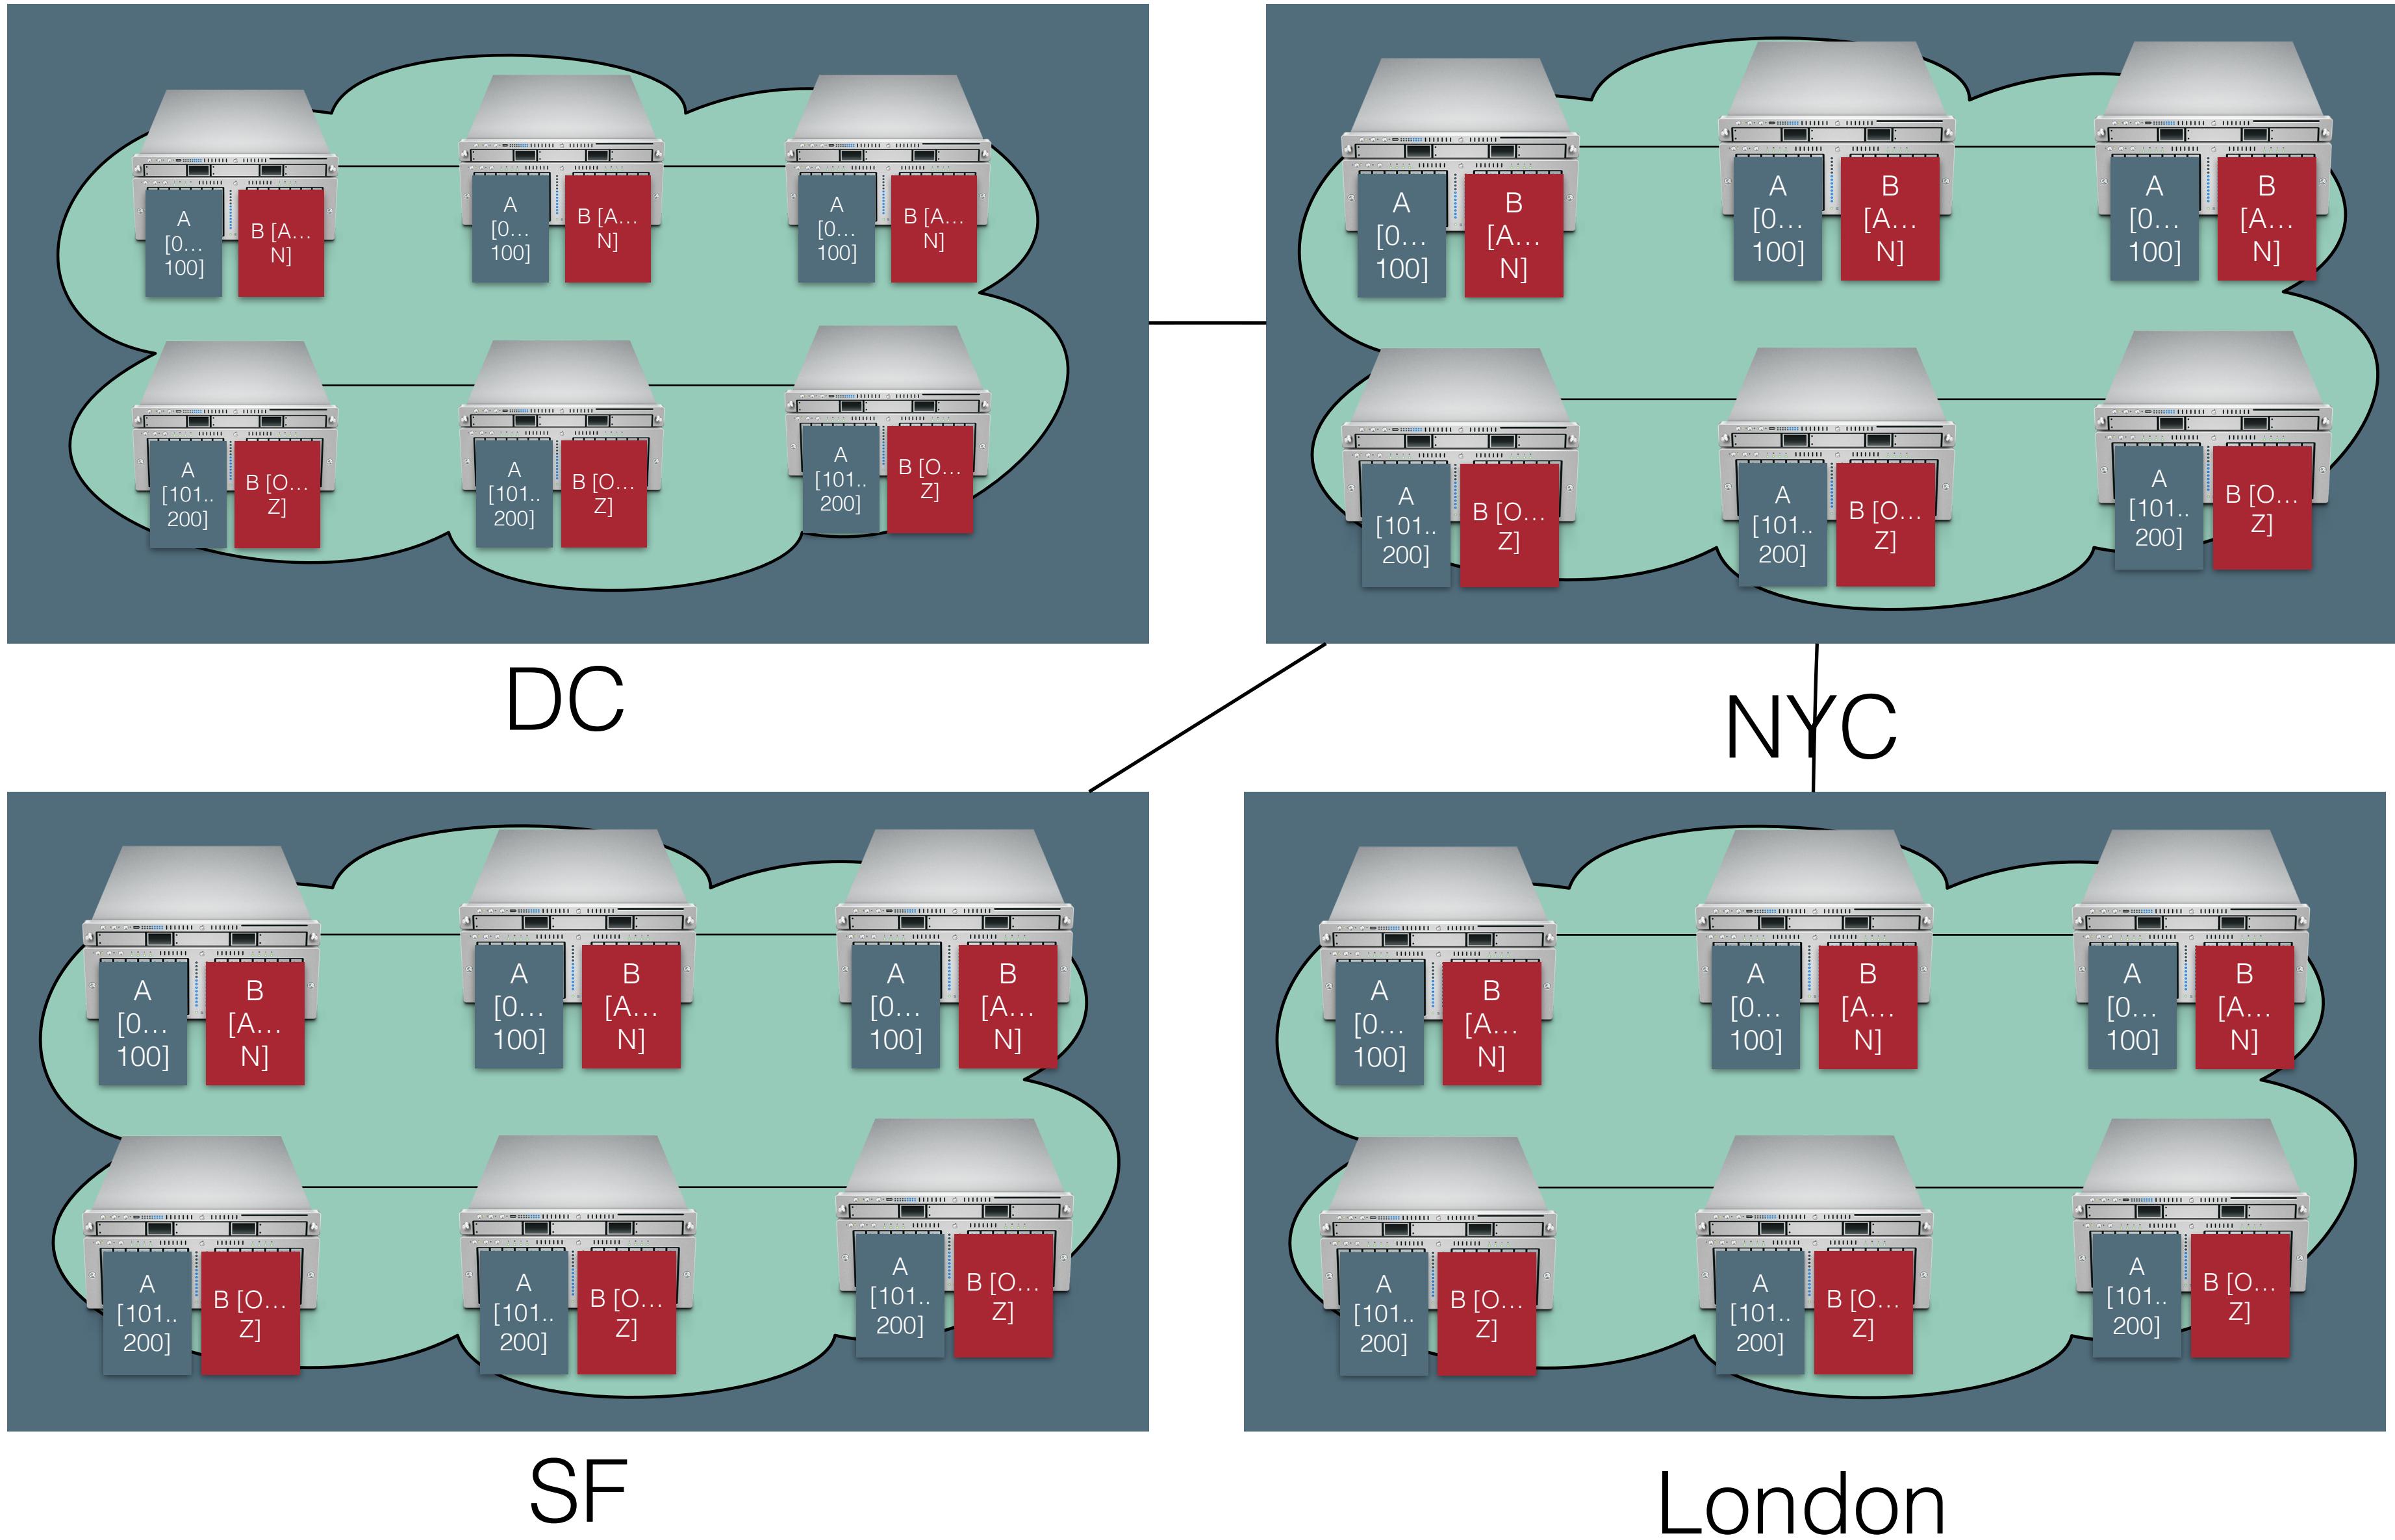
# Partitioning + Replication



# Partitioning + Replication

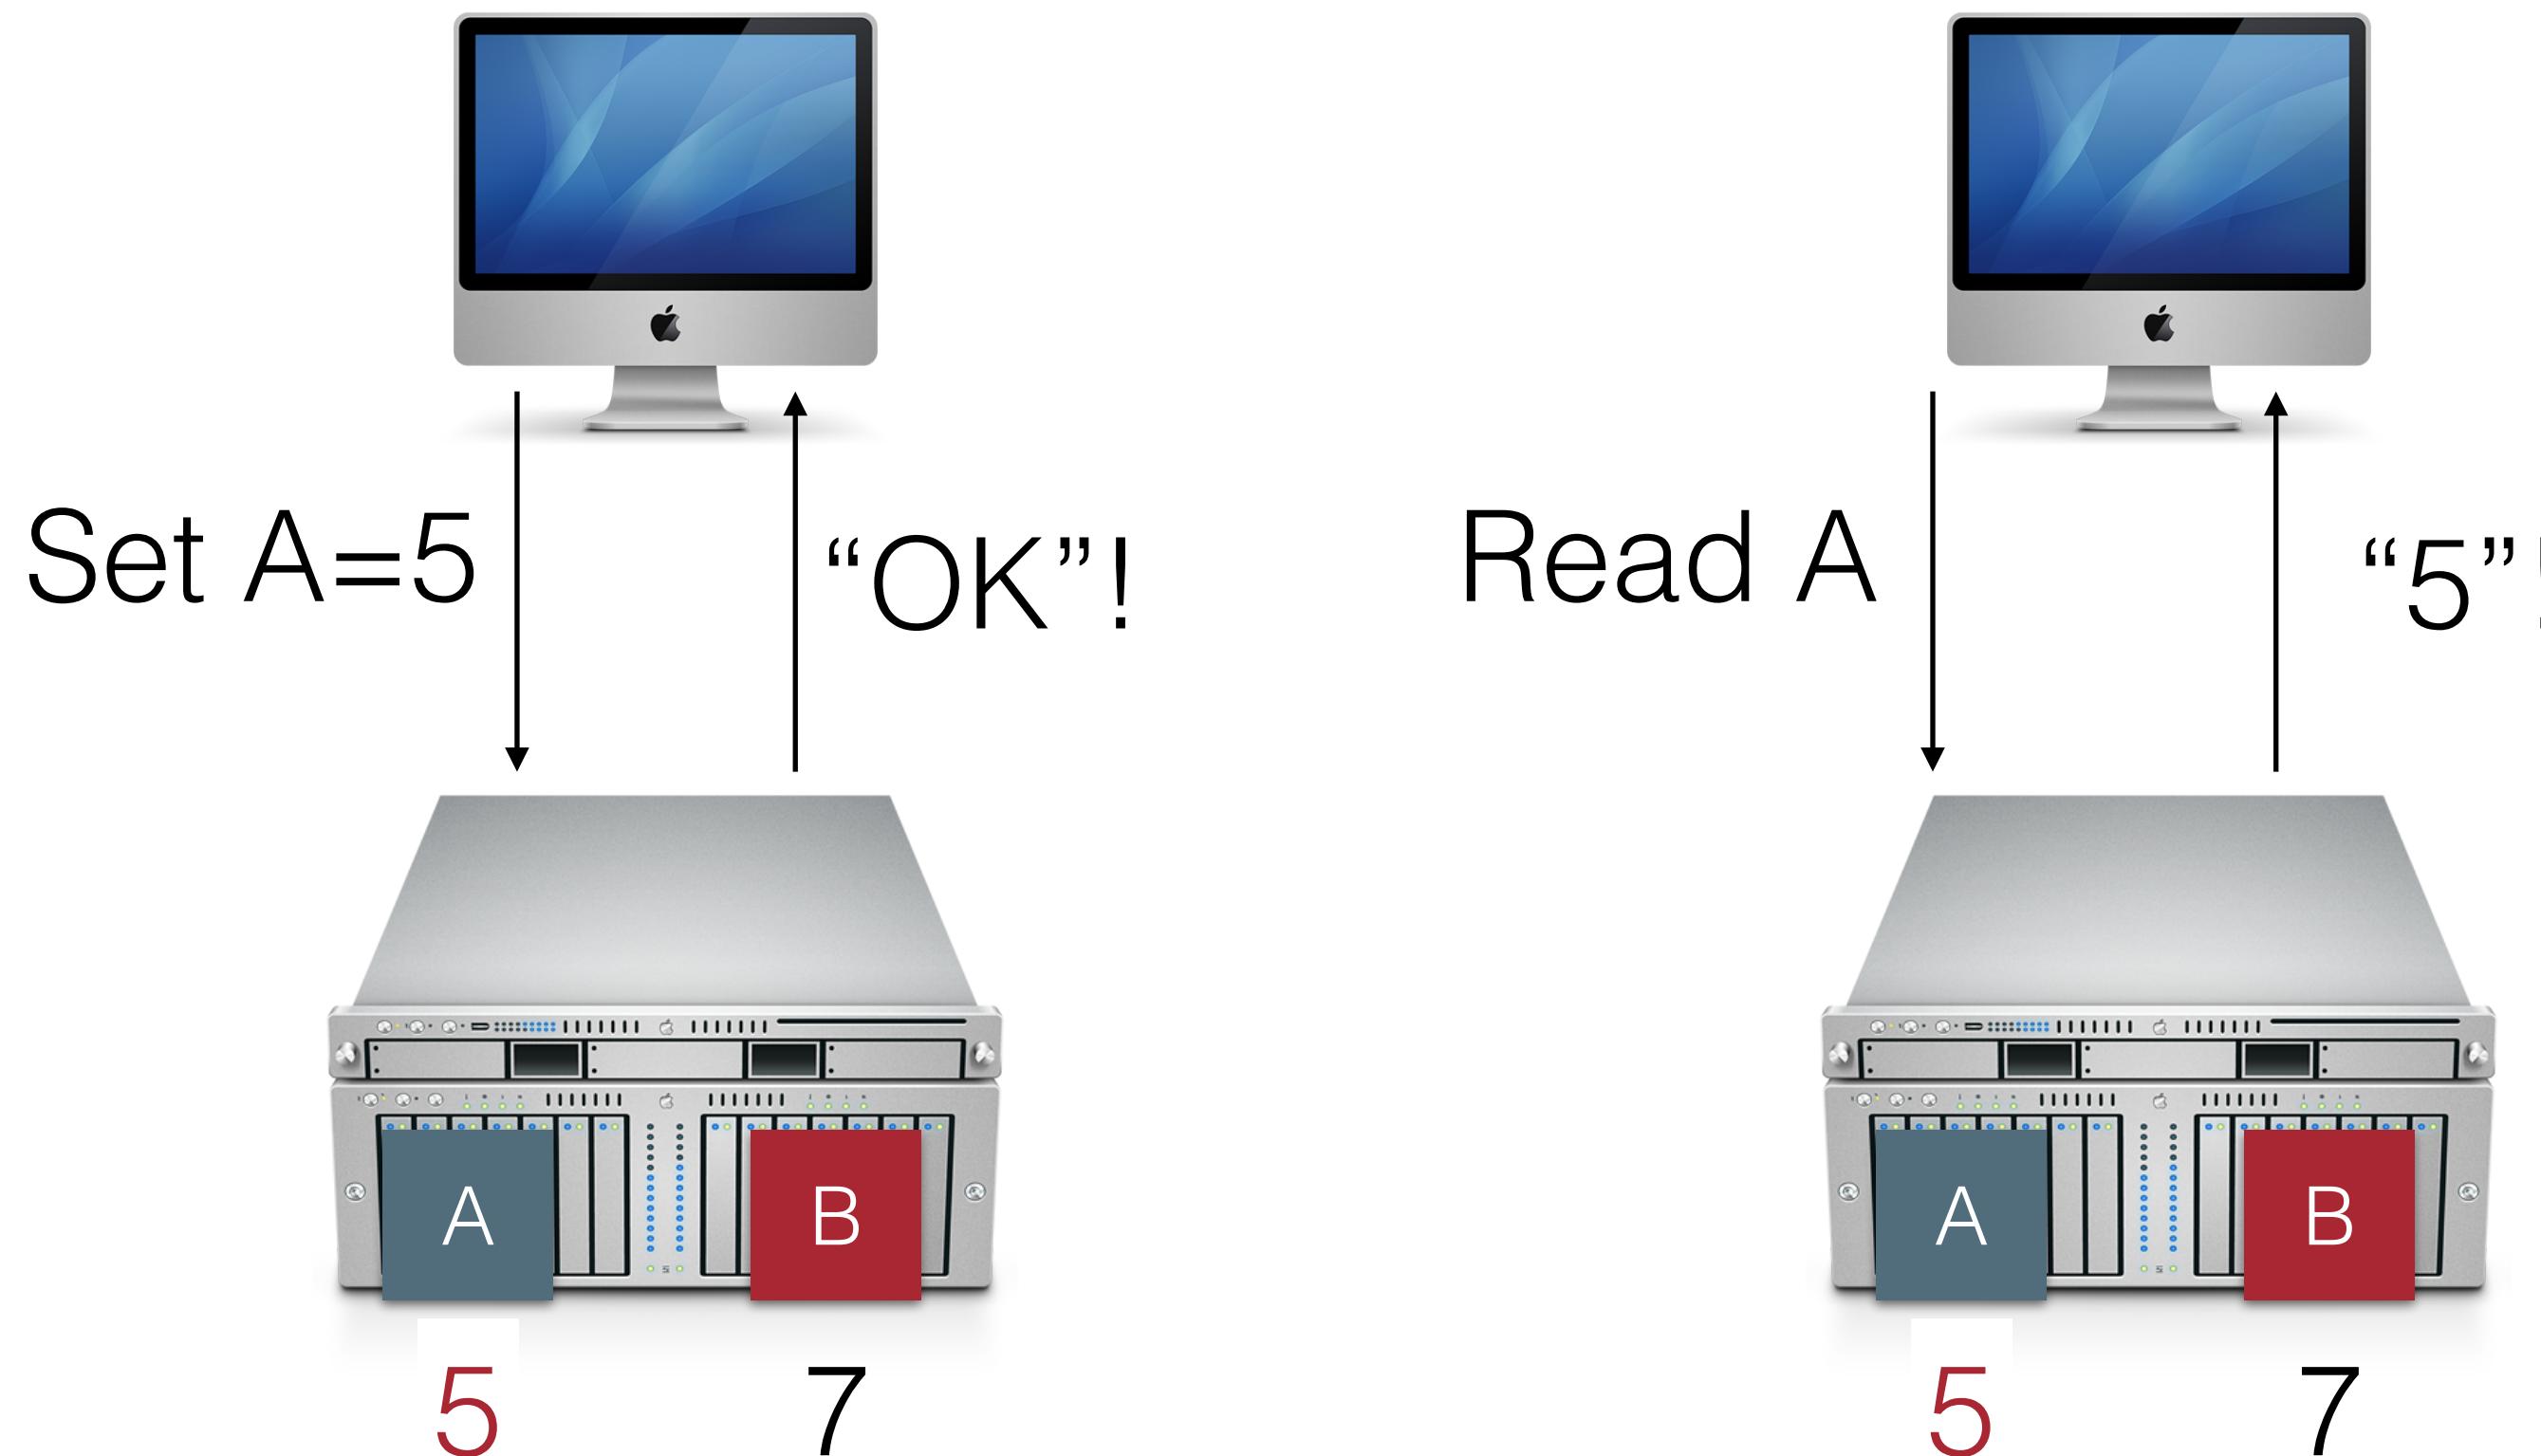


# Partitioning + Replication



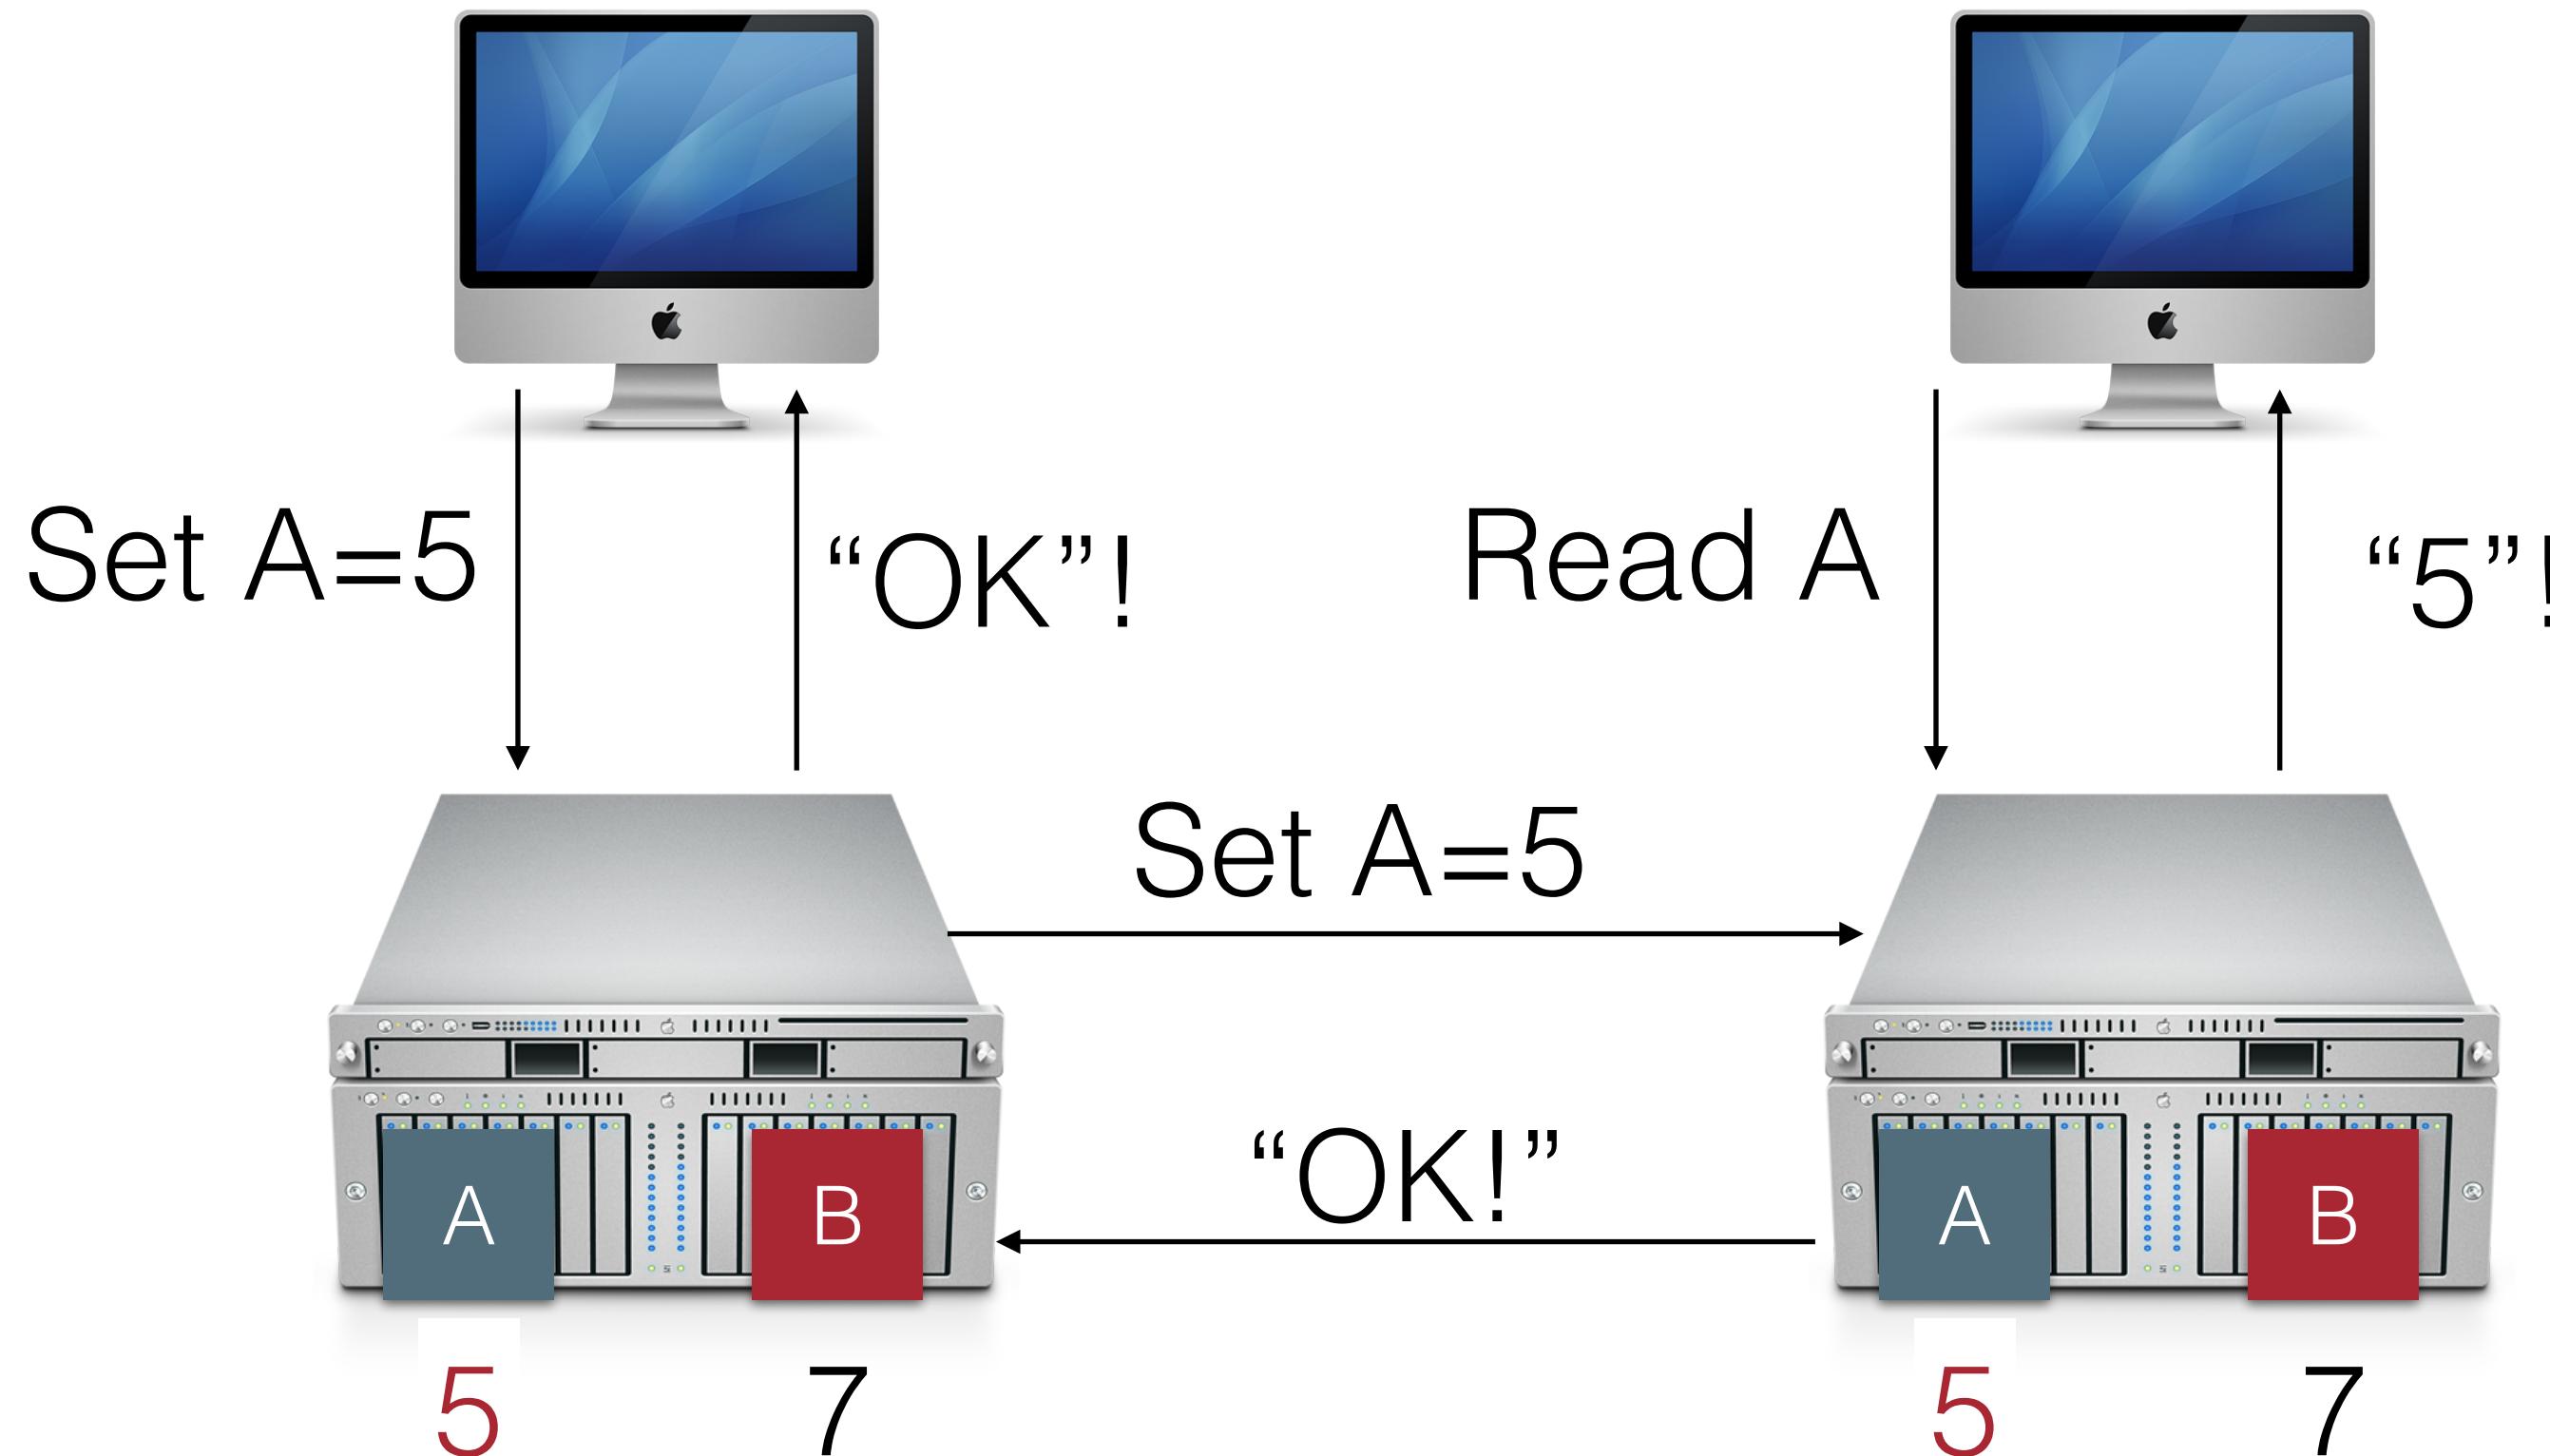
# Replication Problem: Consistency

We probably want our system to work like this



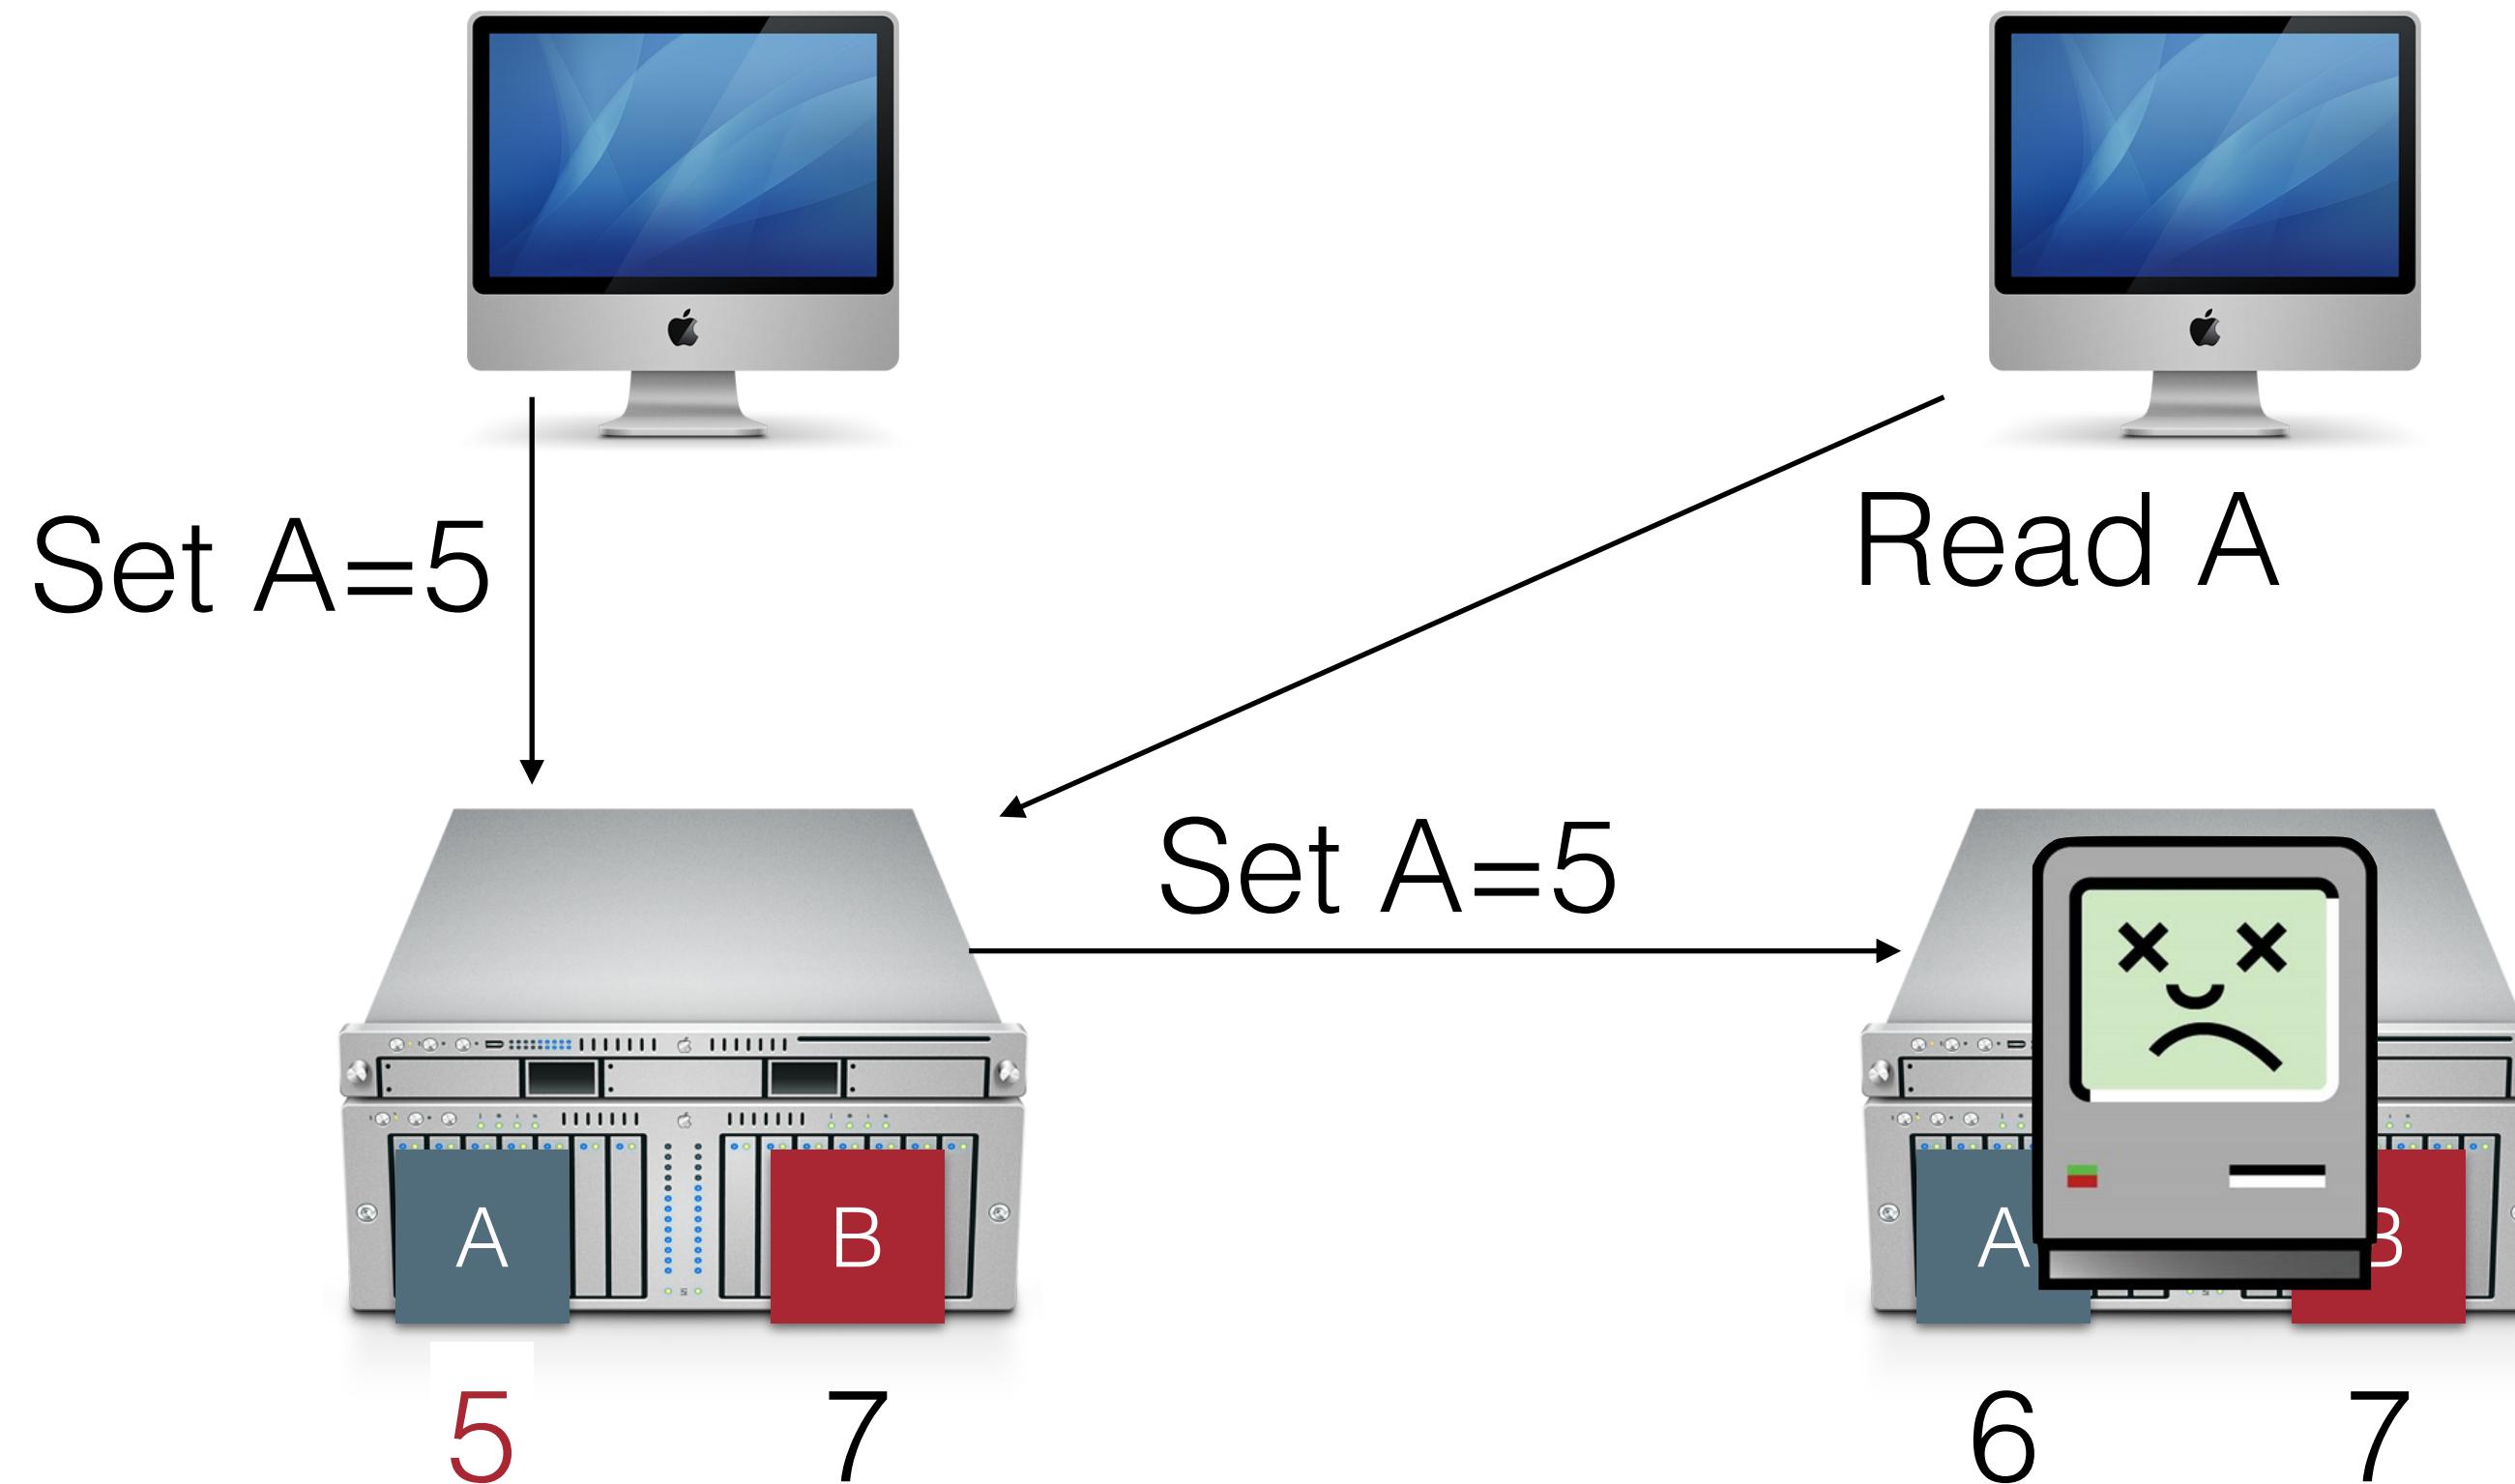
# Sequential Consistency

AKA: Behaves like a single machine would



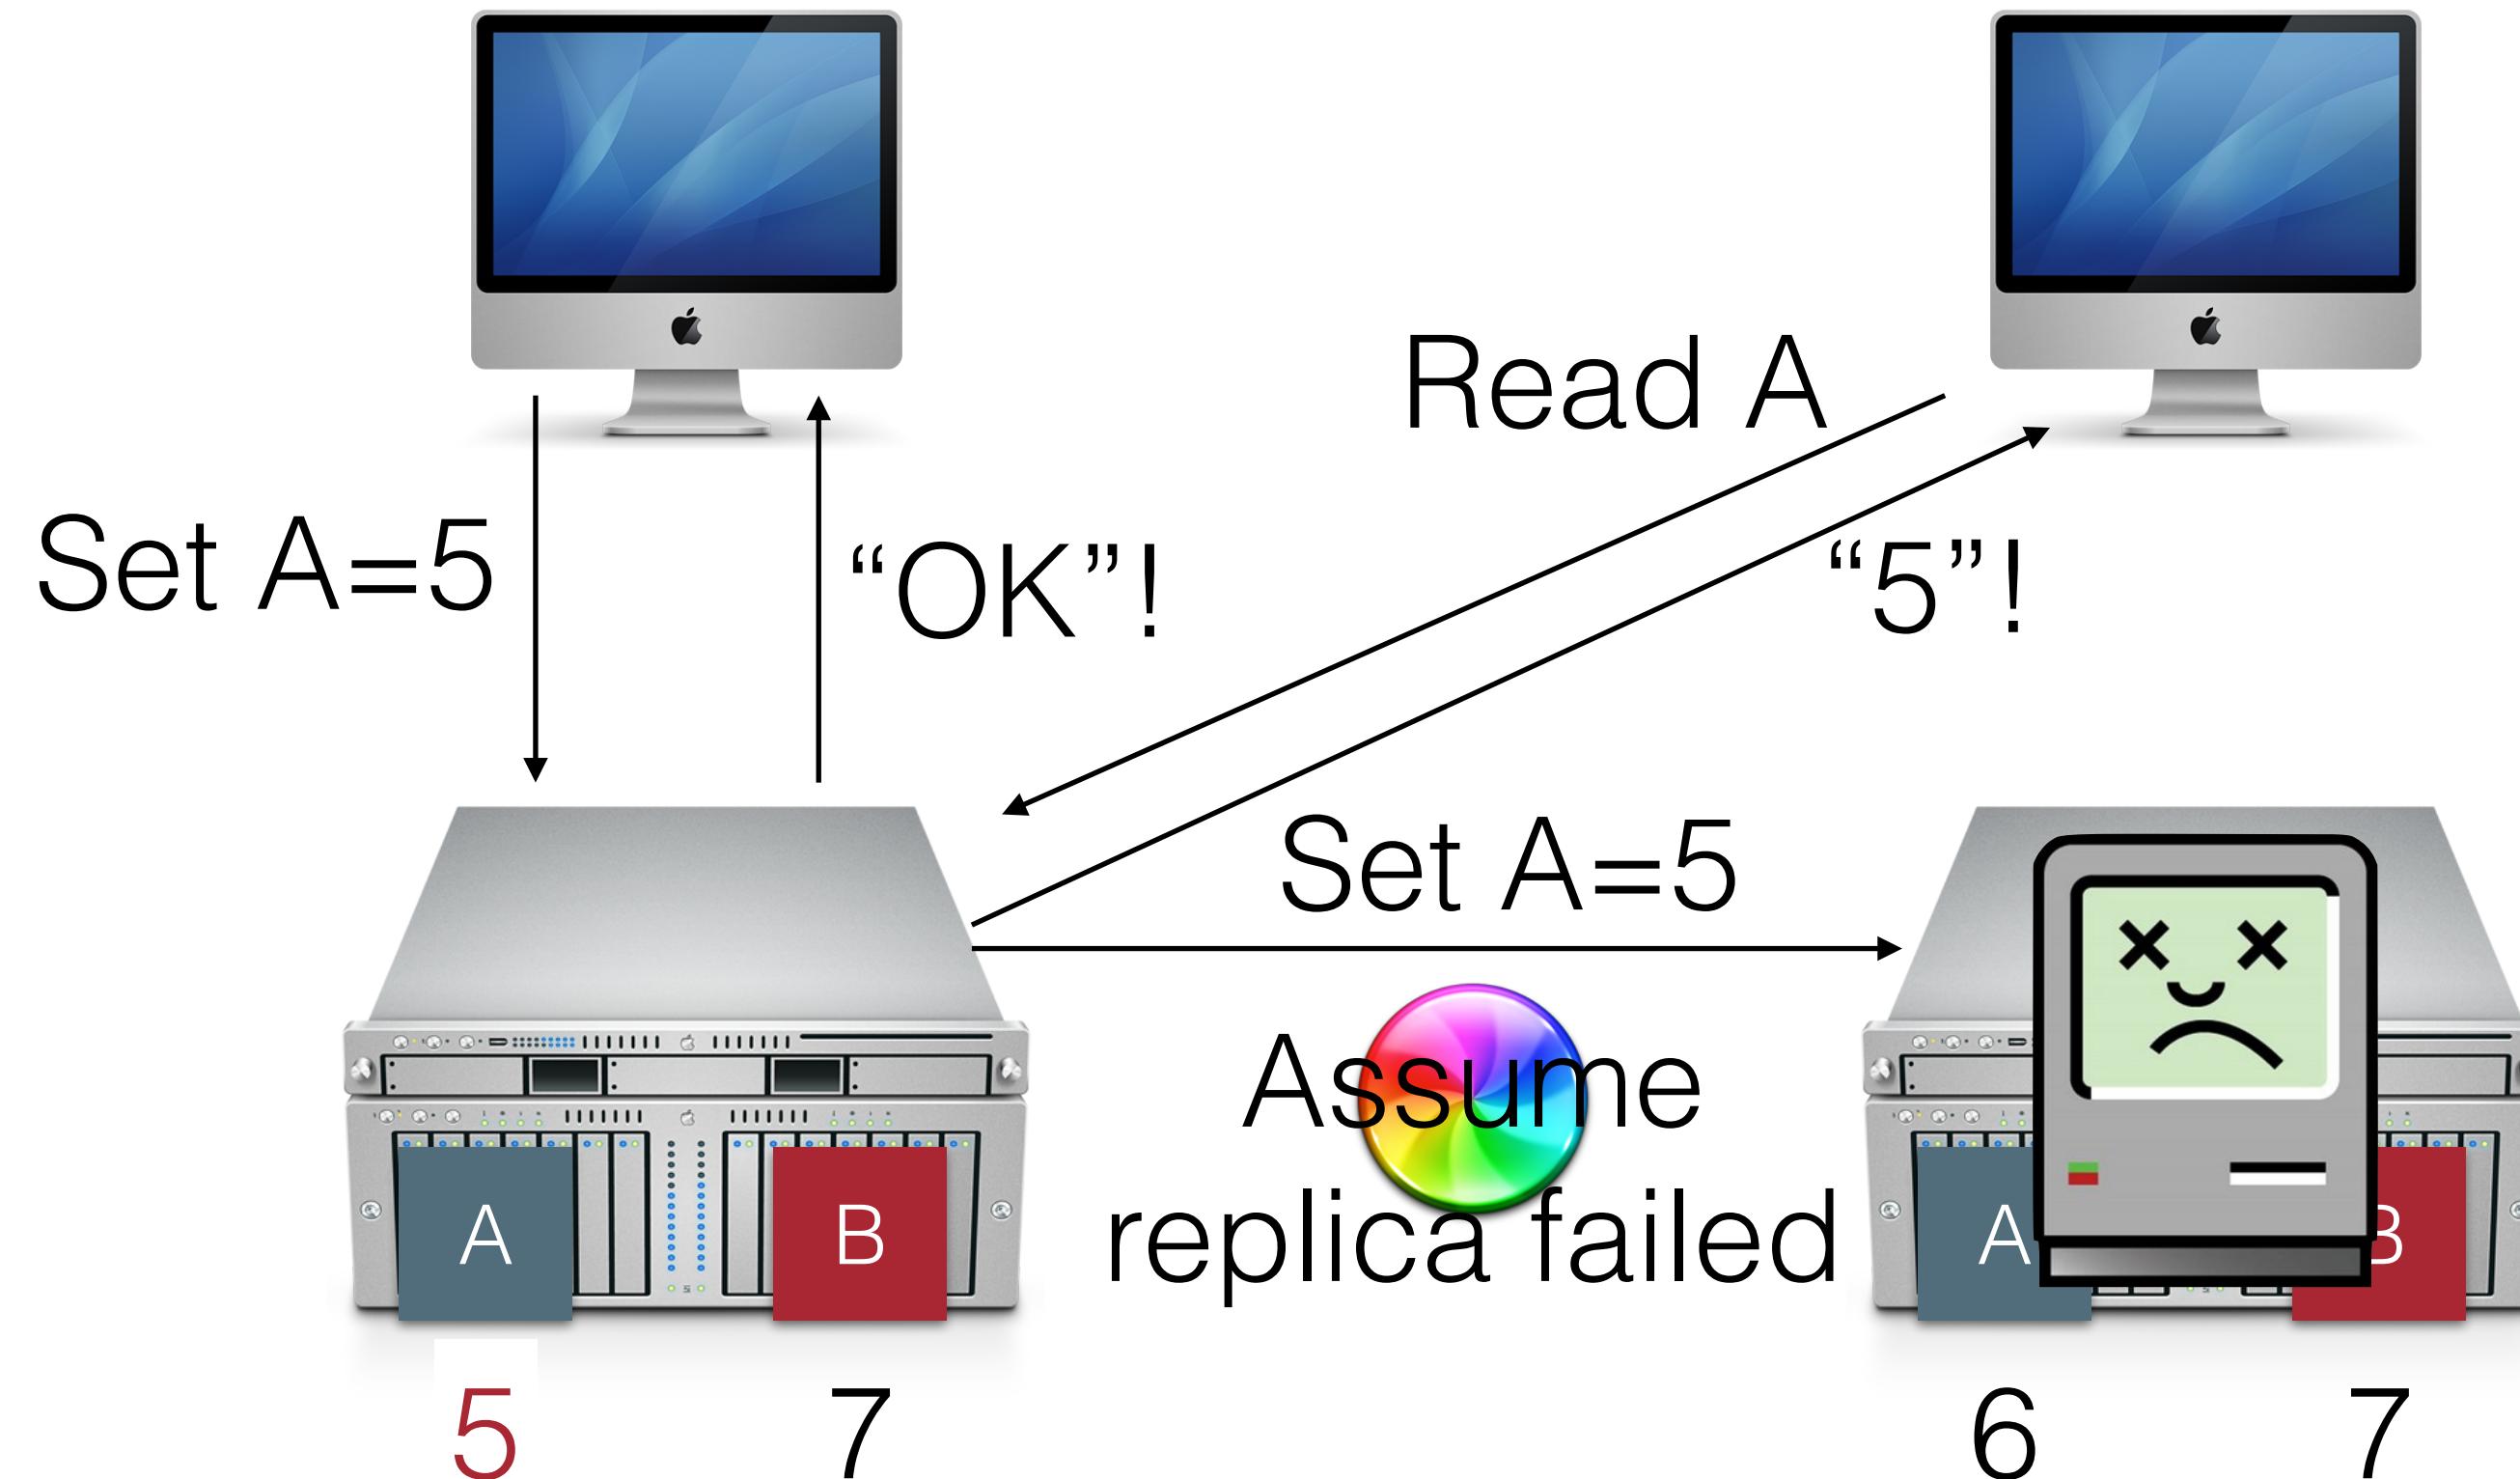
# Availability

If at least one node is online, can we still answer a request?



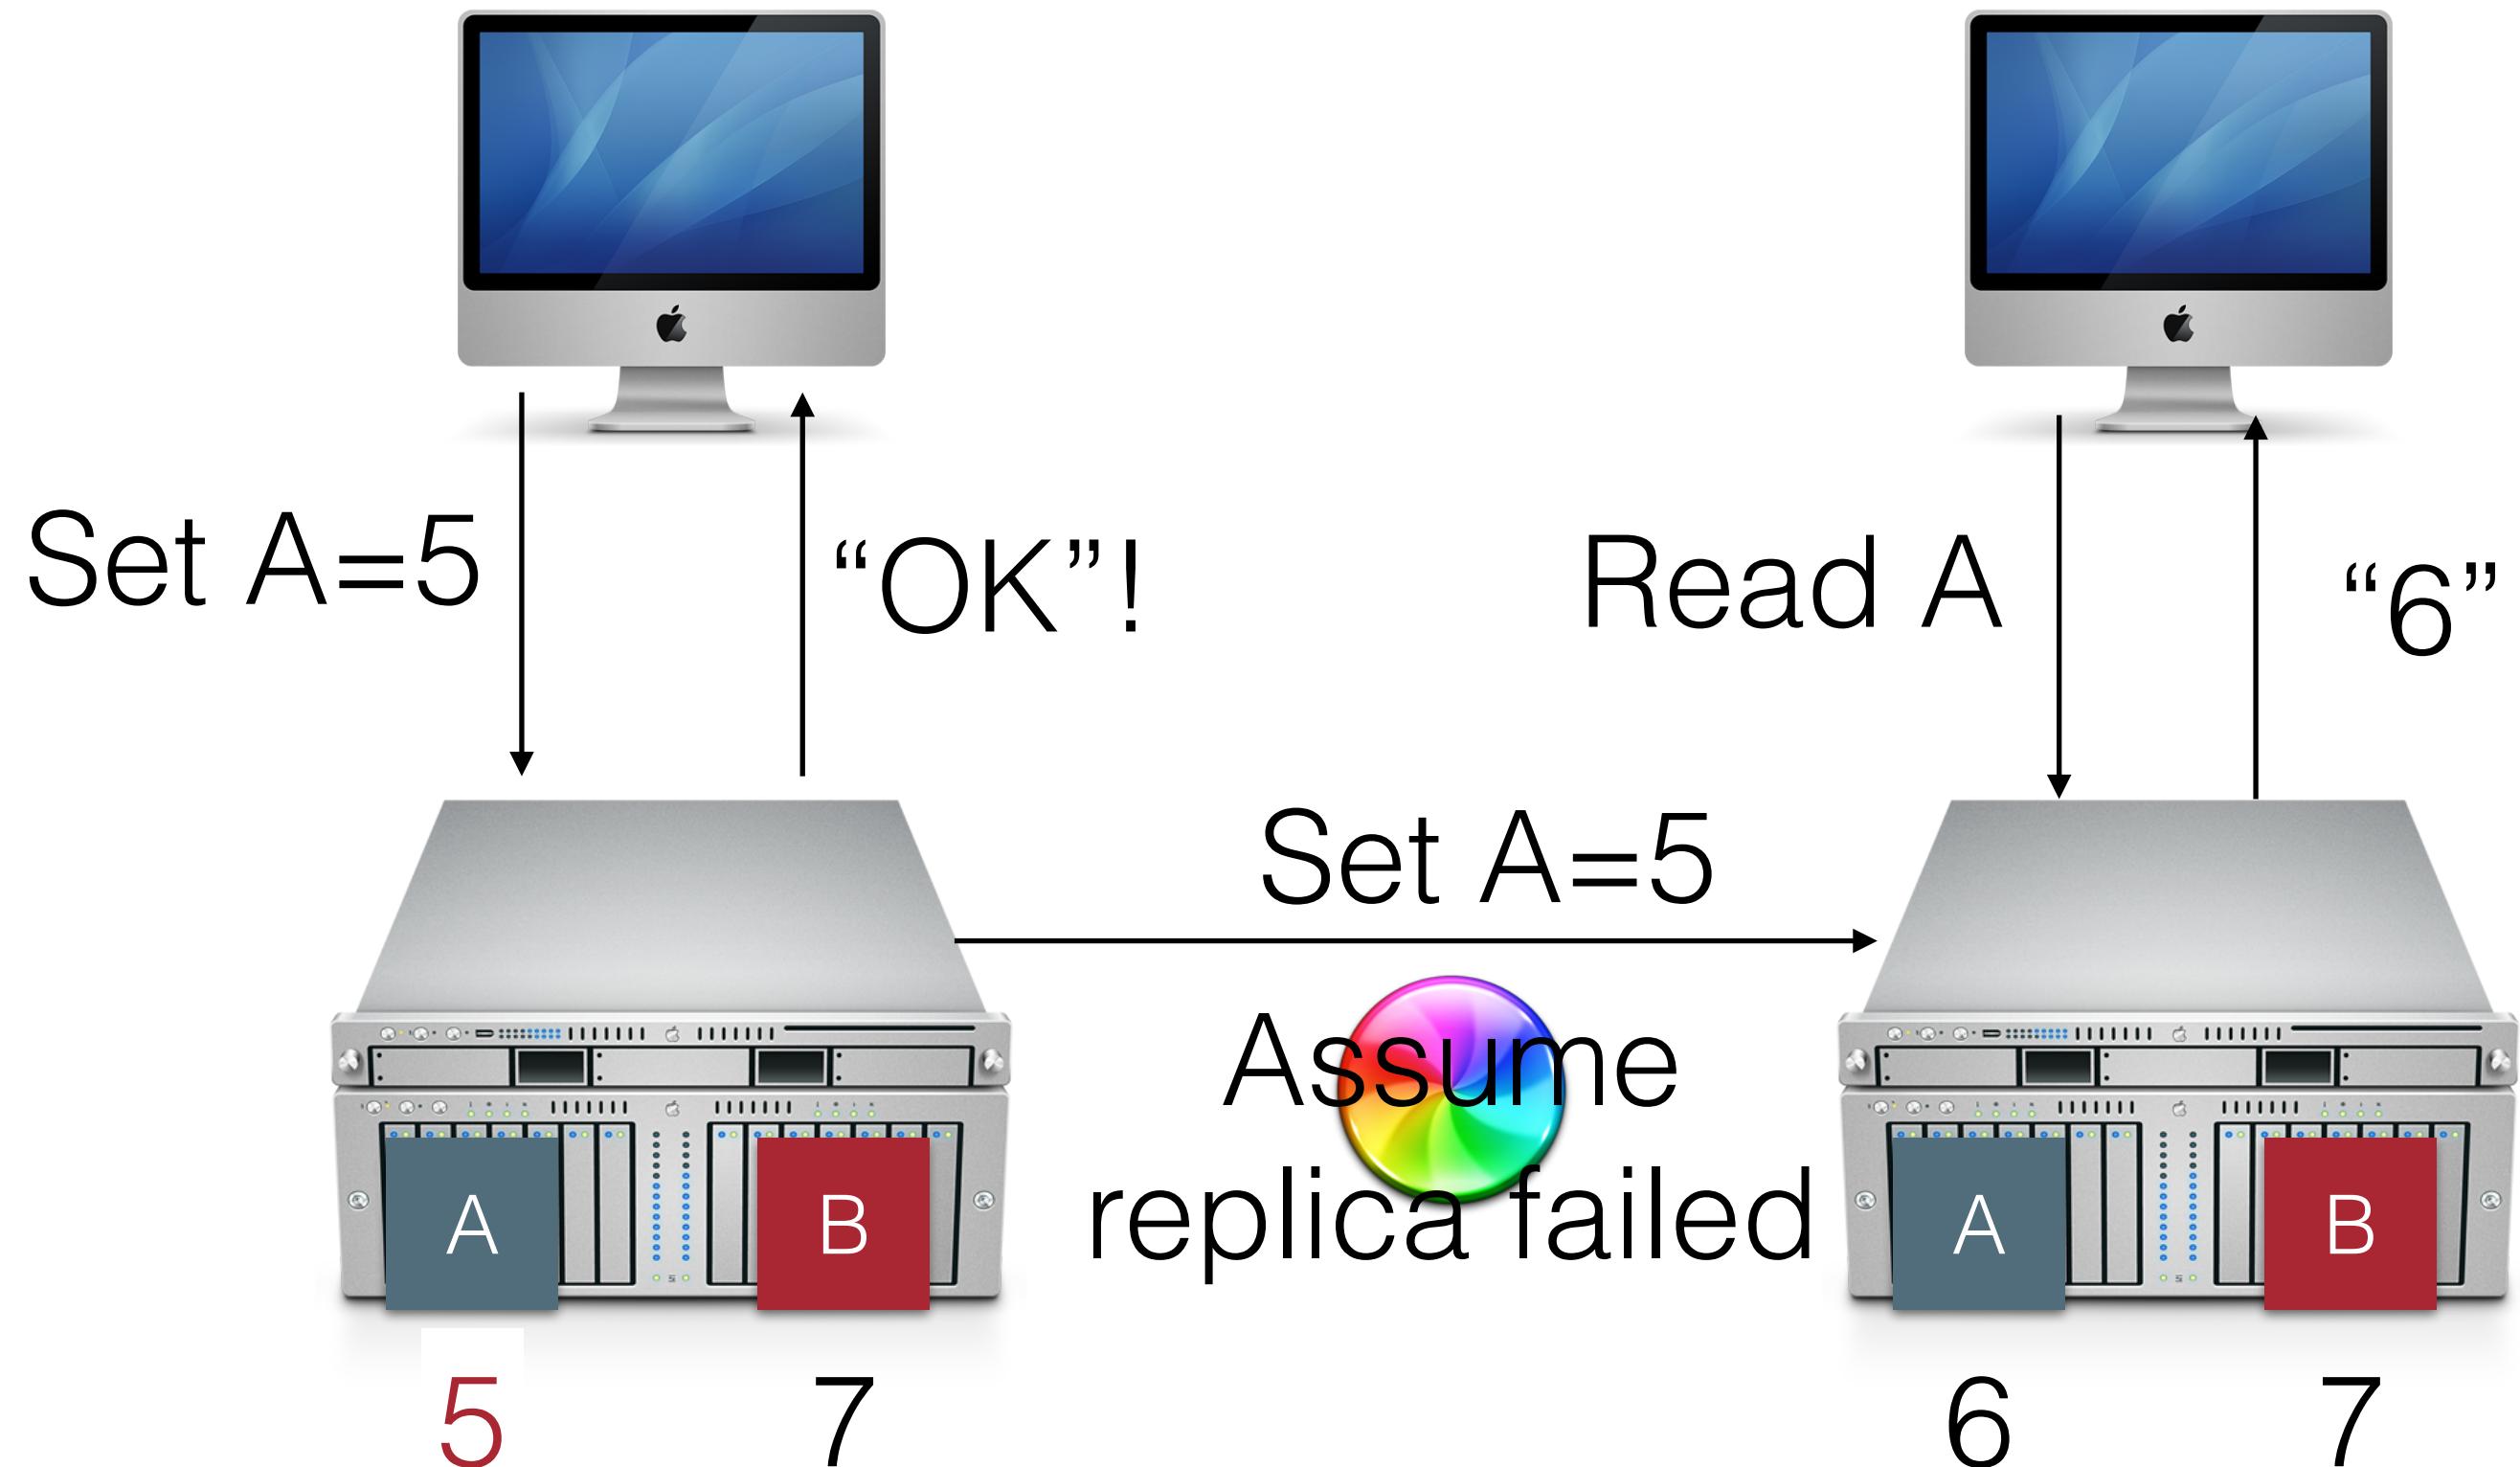
# Consistent + Available

On timeout, assume node is crashed



# Recall: Fallacies of Distributed Systems

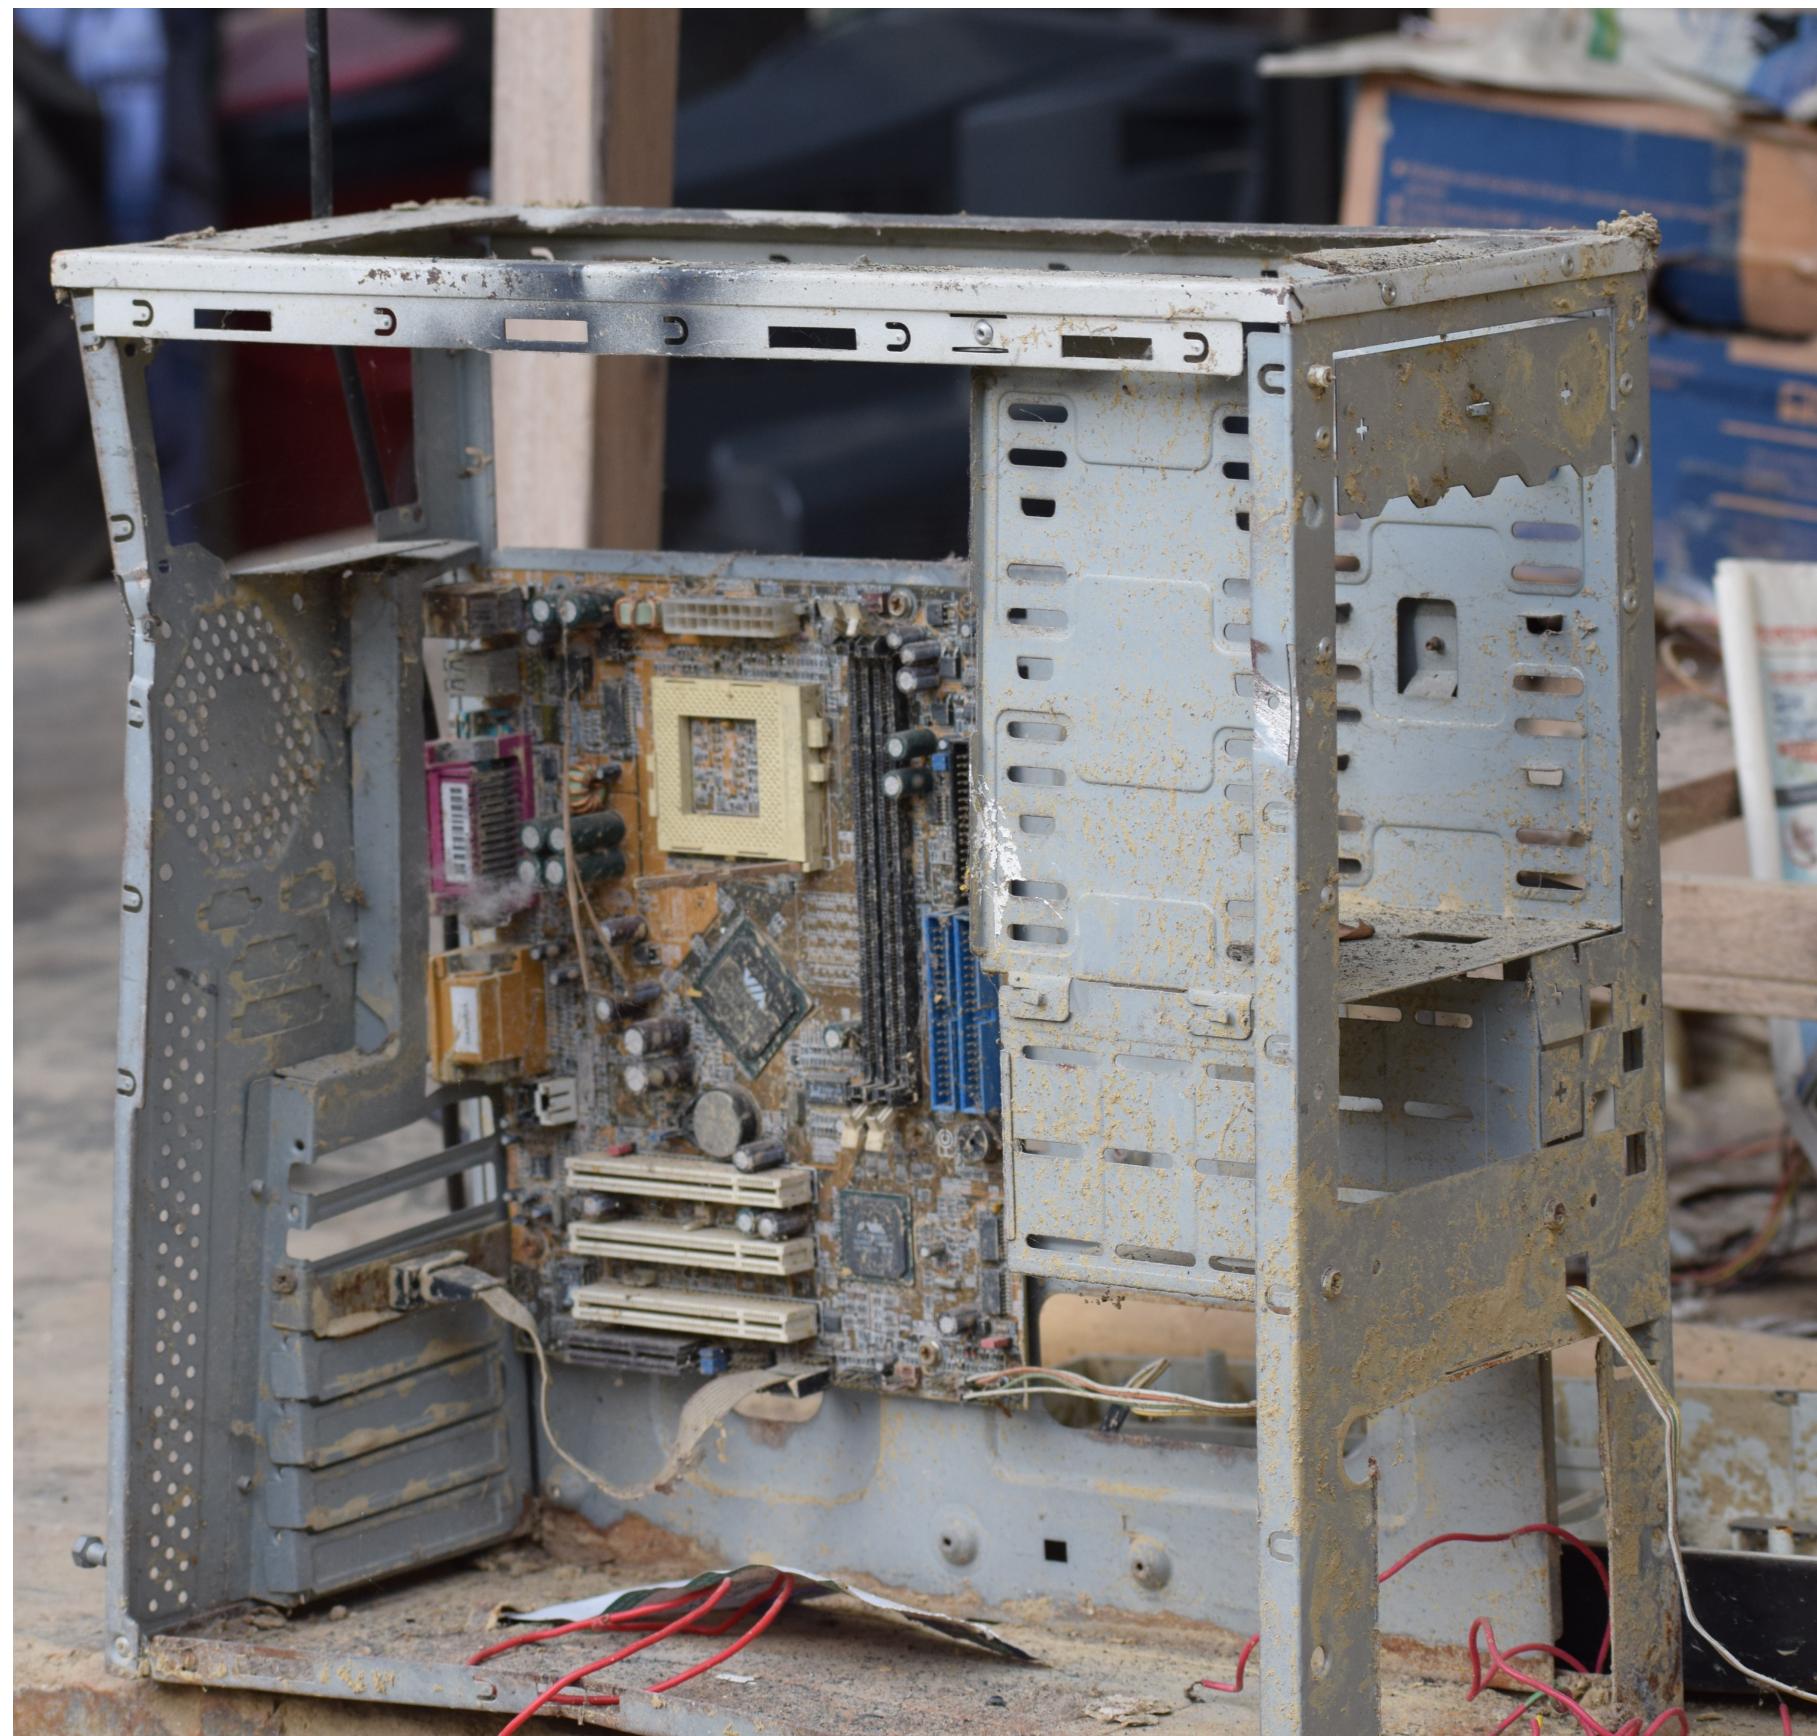
What if *the network fails*?



# Shared Fate

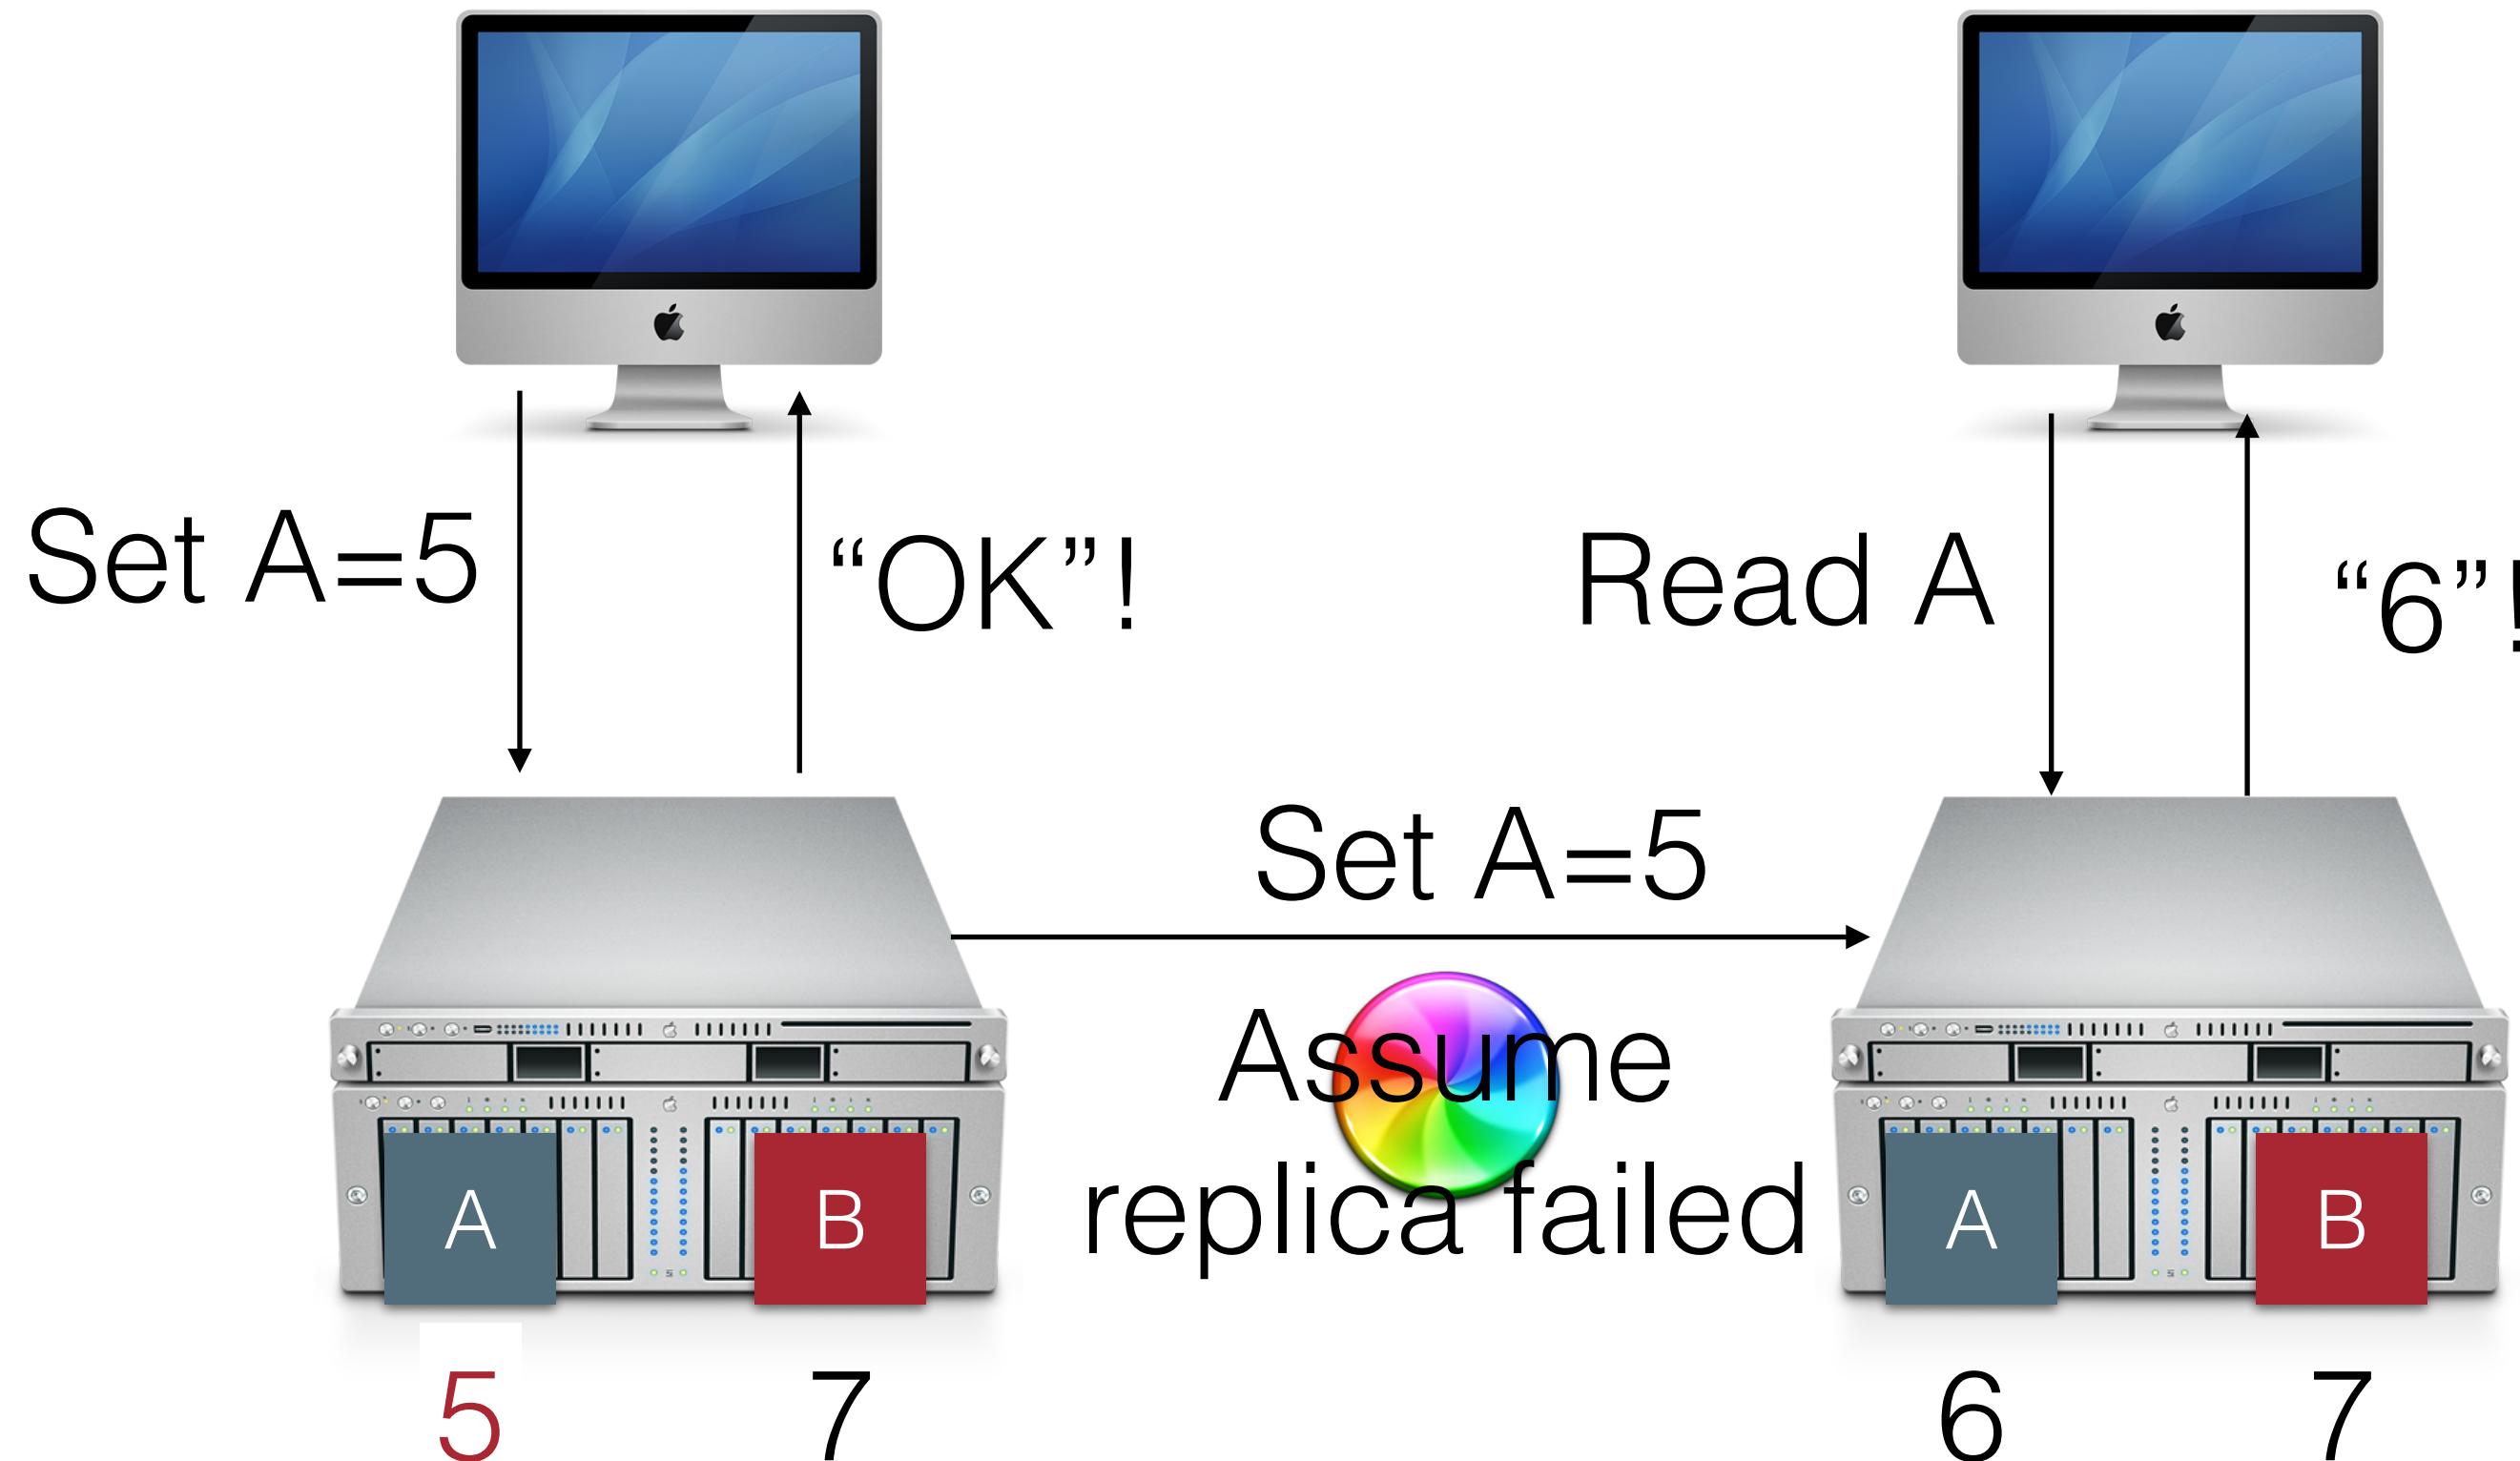
Are you still there?

- Two methods/threads/processes running on the same computer generally have **shared fate** [Crashed/not]
- When two machines in a distributed system can't talk to each other, how do we know if the other is crashed?
- We call this a **split brain** problem



# Network Partitions

The network might fail *arbitrarily*!

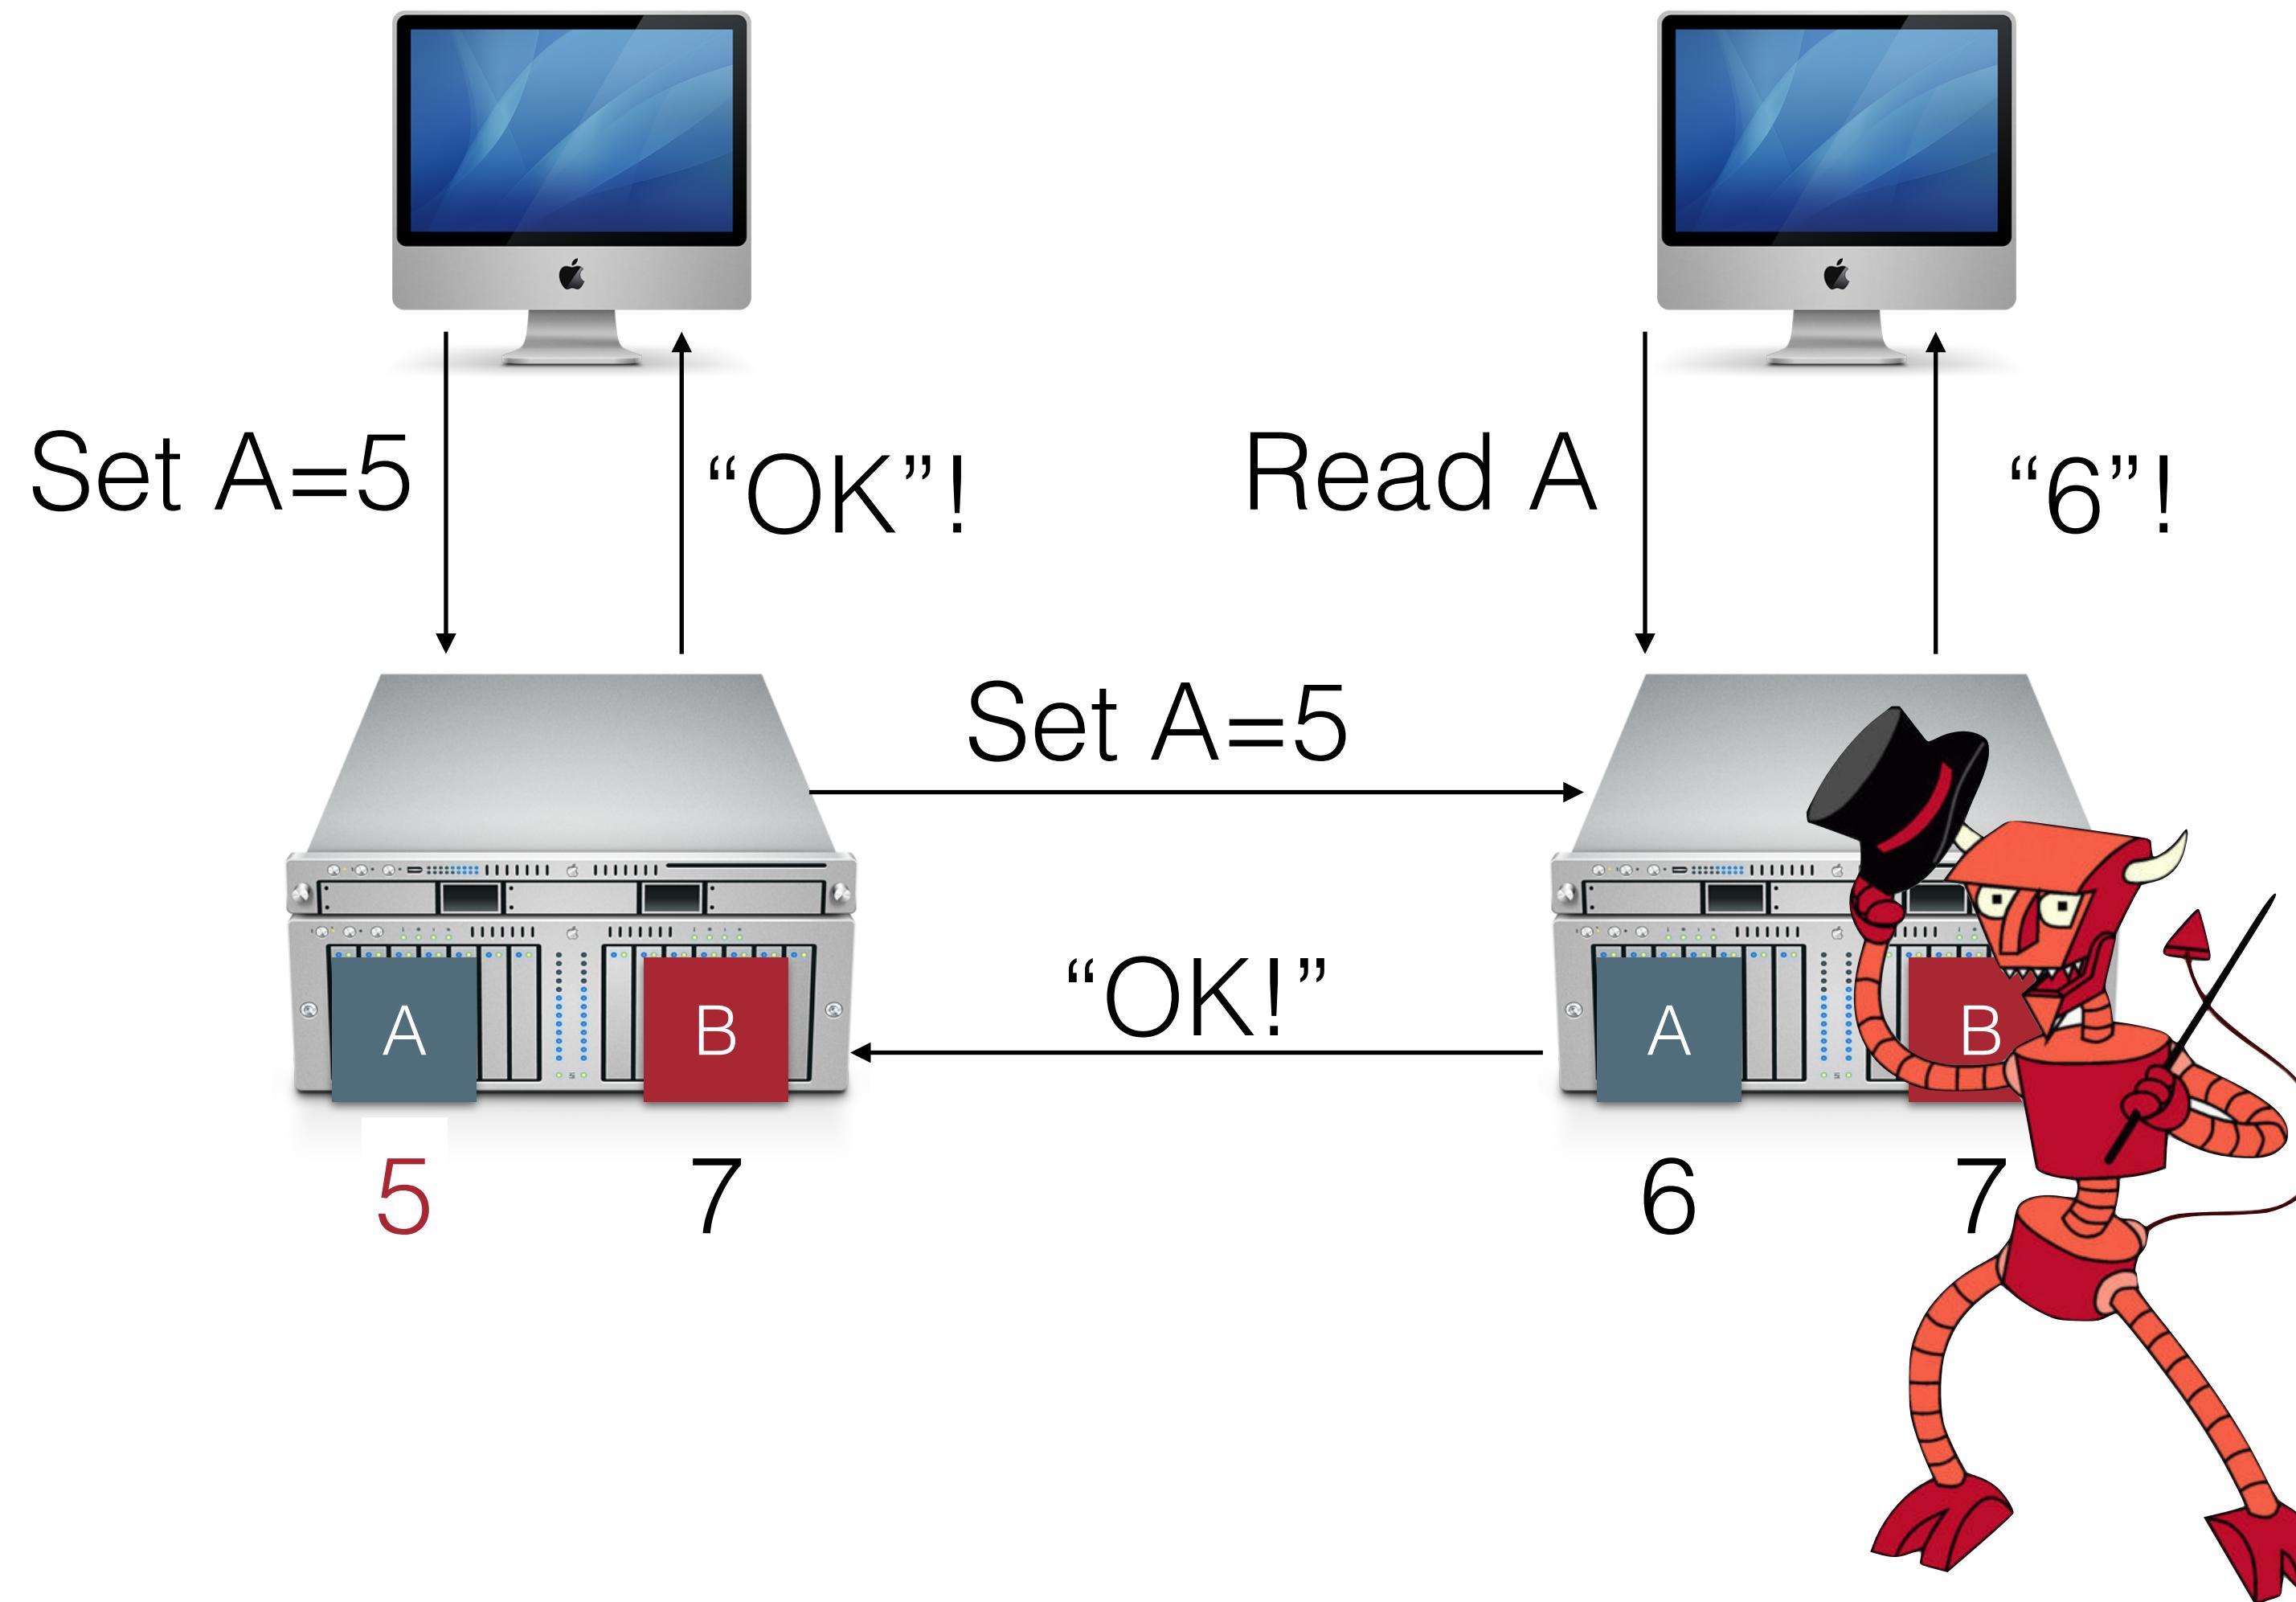


# CAP Theorem

- Pick two of three:
  - Consistency: All nodes see the same data at the same time (strong consistency)
  - Availability: Individual node failures do not prevent survivors from continuing to operate
  - Partition tolerance: The system continues to operate despite message loss (from network and/or node failure)

# Byzantine Faults

Unfortunately, still more things can go wrong



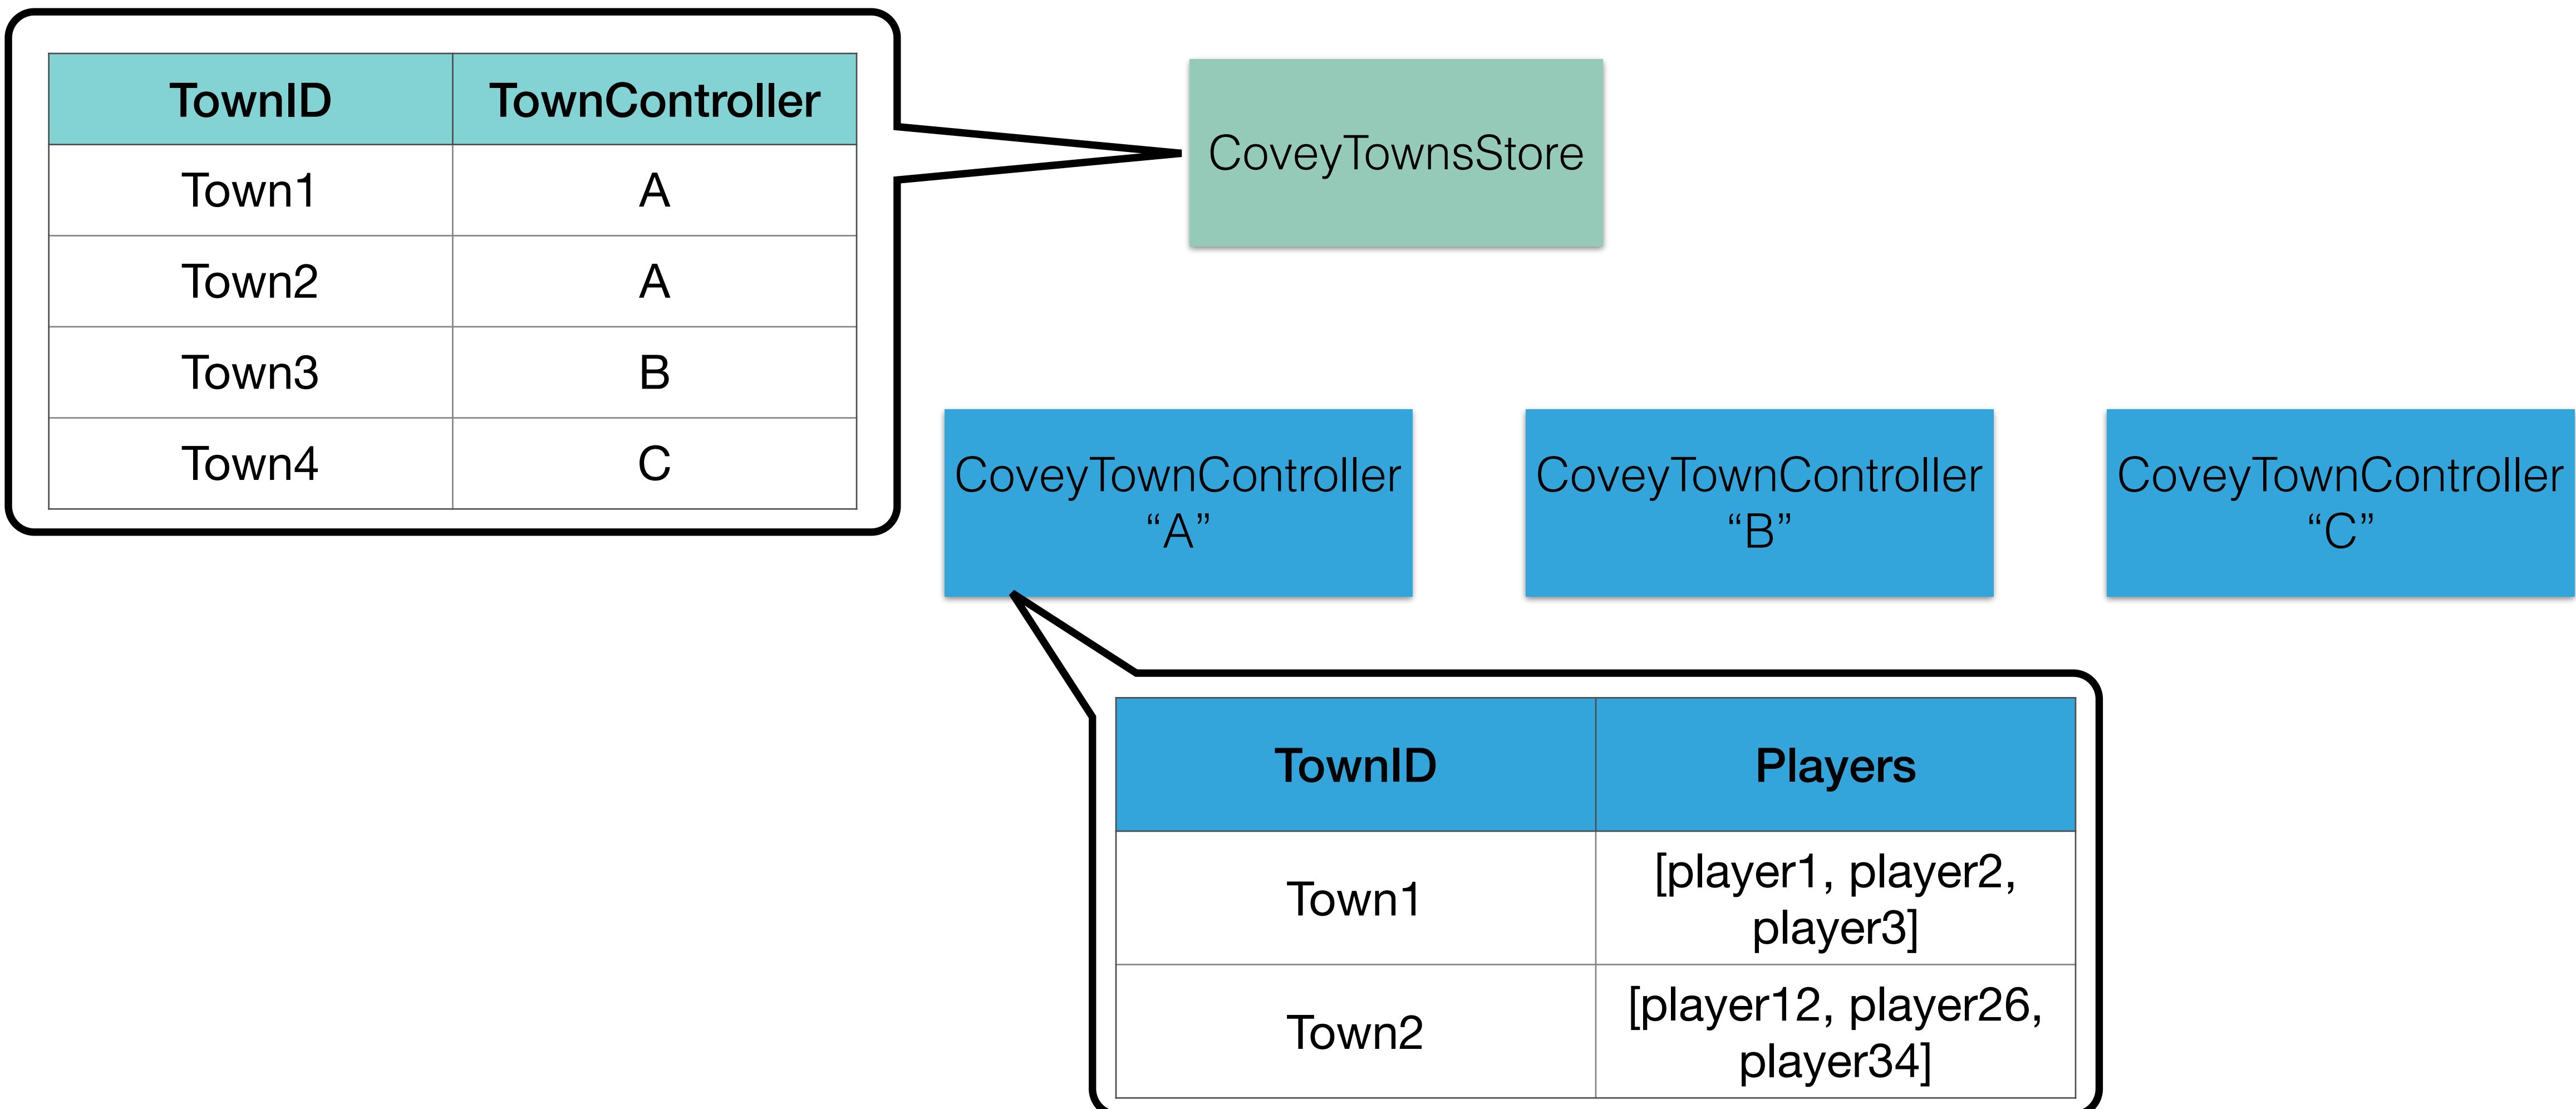
# Distributed Software Engineering Abstractions

## Key Question: Consistency vs Availability

- Distributed system will never match exact semantics of non-distributed system
- For replication do we value more: guaranteed consistency (looks like a single machine) or guaranteed availability (sometimes read stale data)?
  - For a lock server?
  - For the order of tweets on twitter?
- For partitioning: Where can we draw the line?

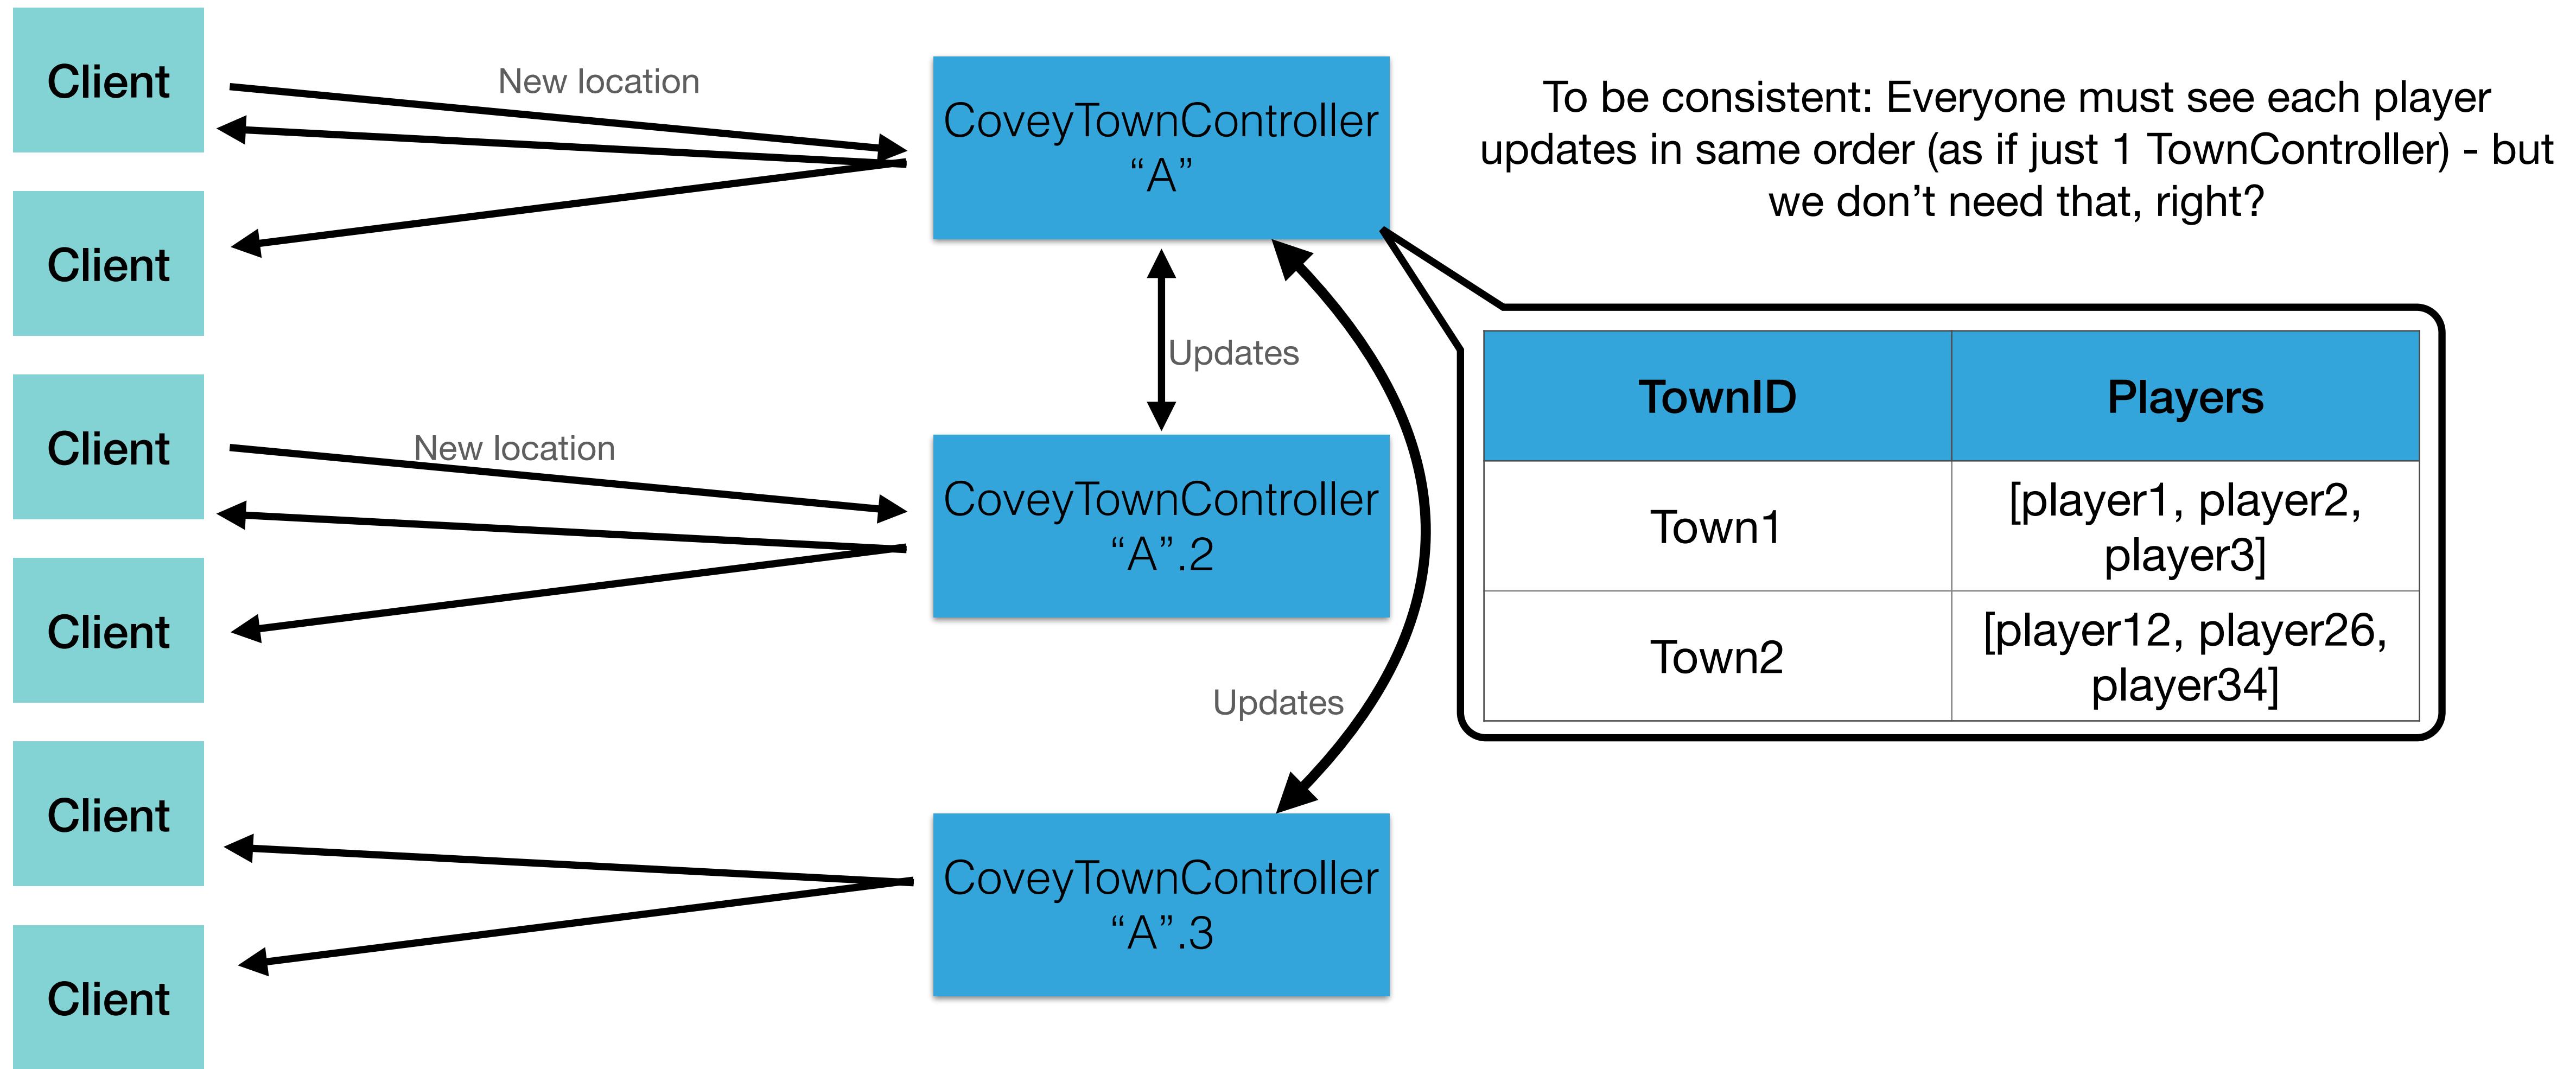
# Distributing Covey.Town

## Partitioning: Separate responsibilities by town



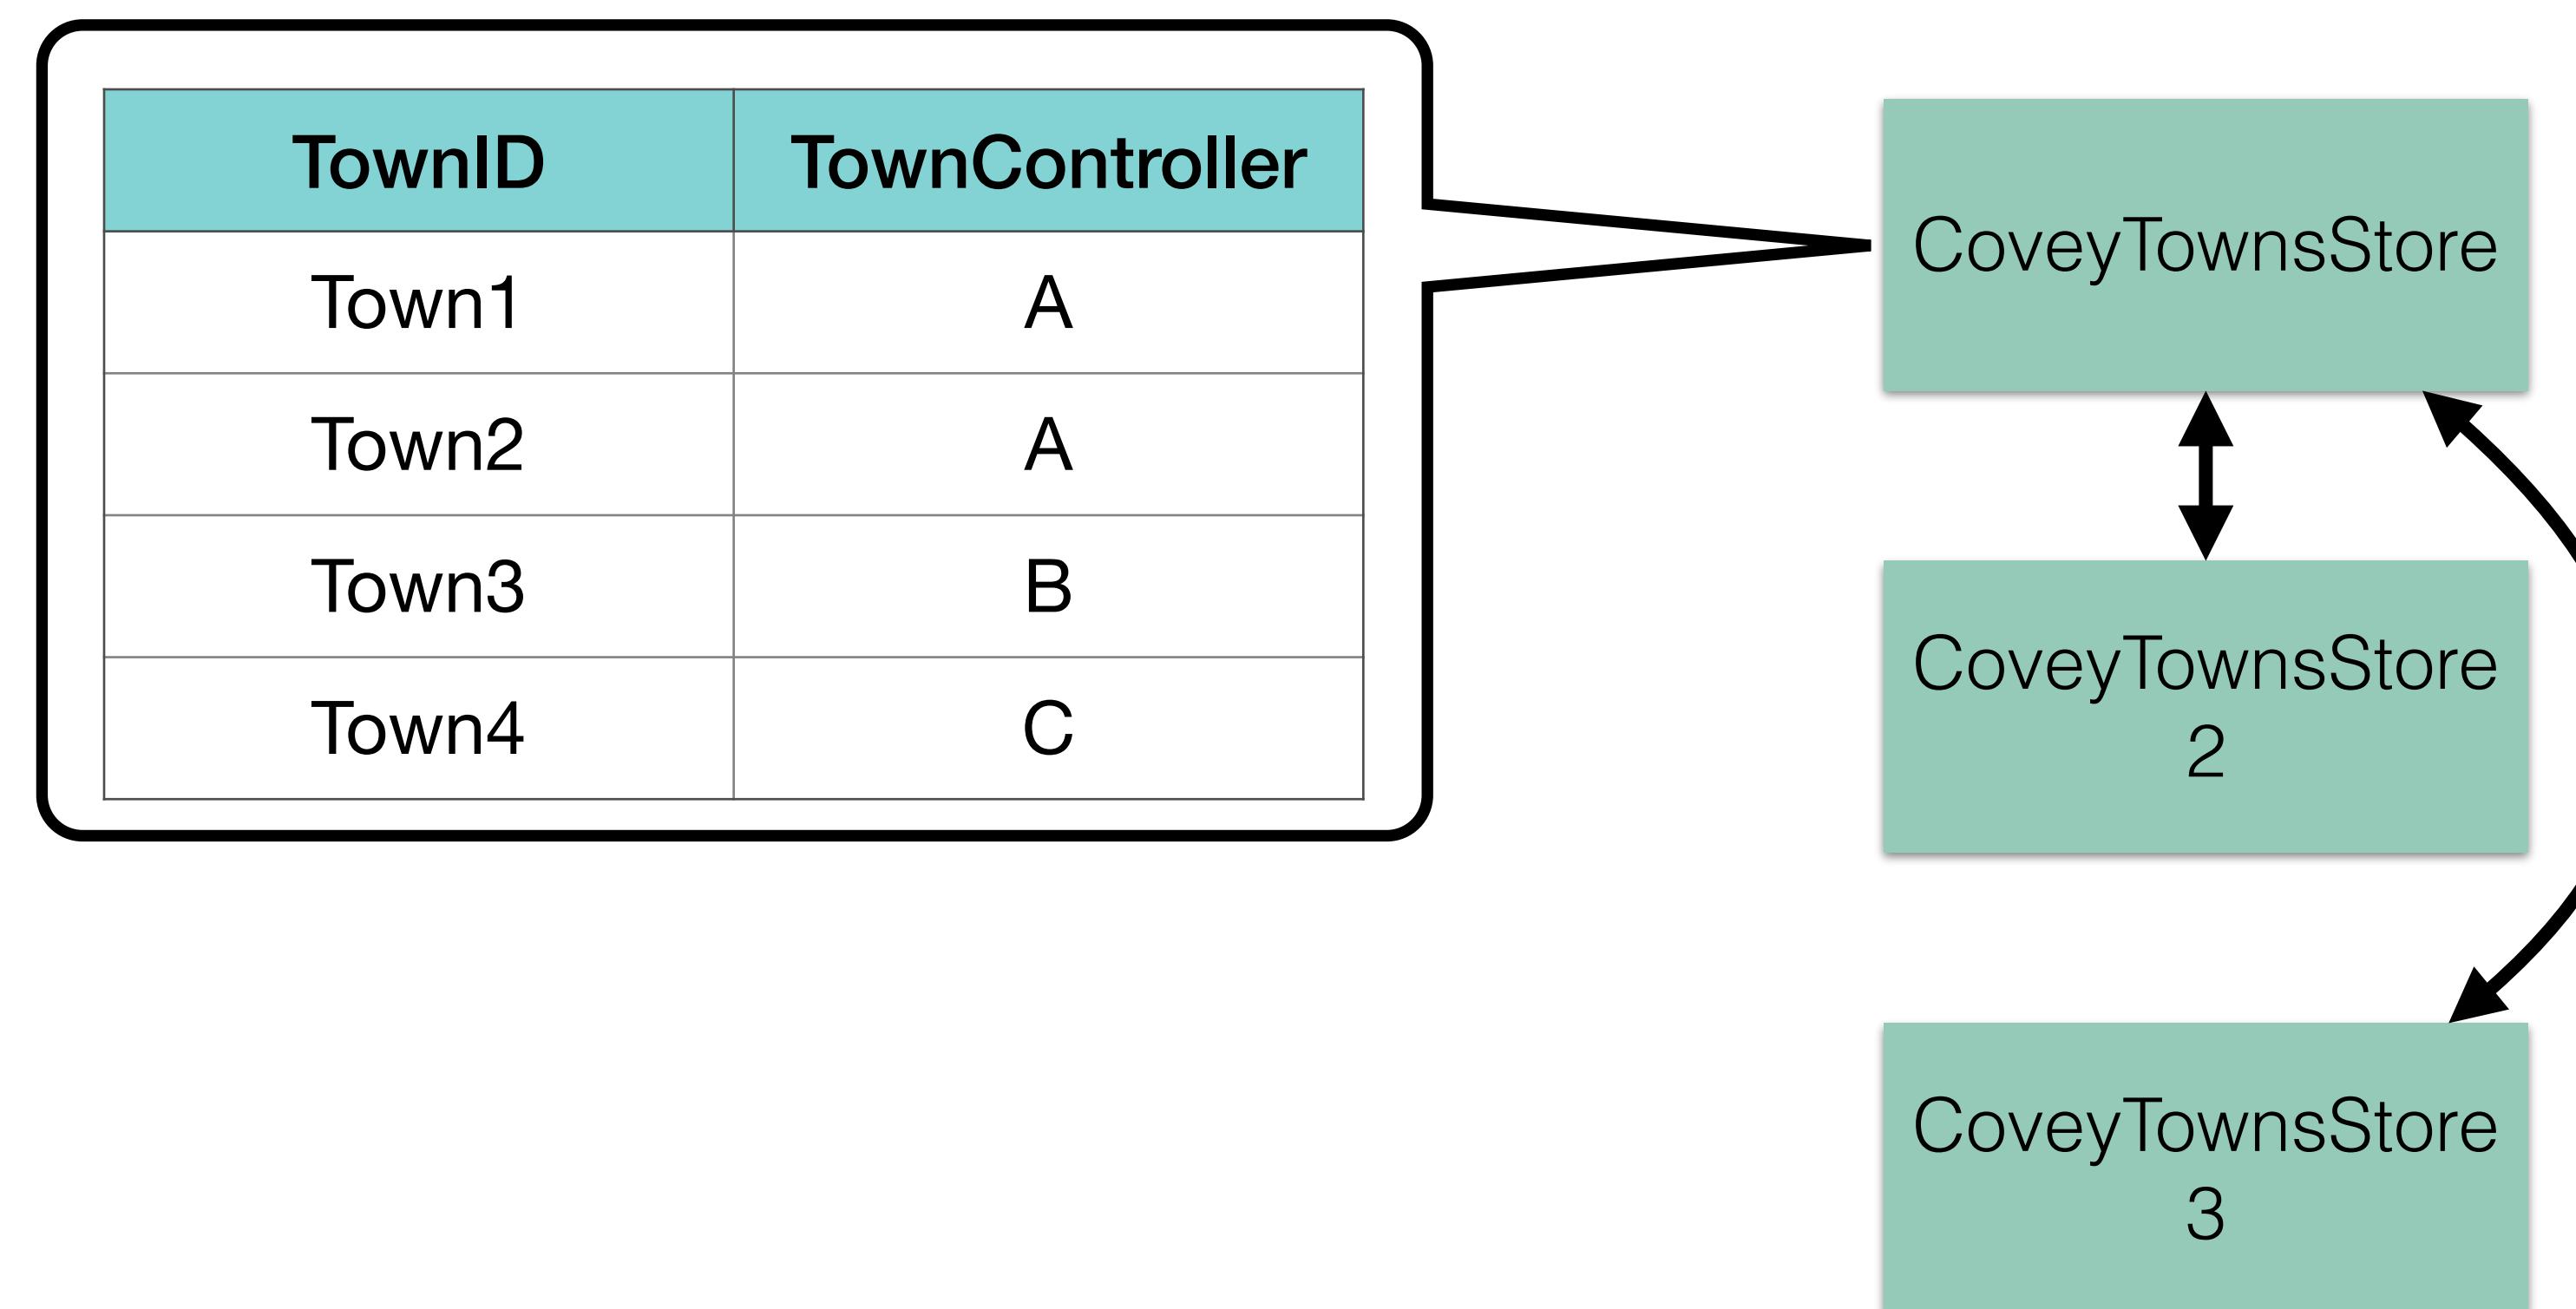
# Distributing Covey.Town

## Replicating the Town Controller



# Distributing Covey.Town

## Replicating the Town Store



What are our consistency challenges?

- Each town must have a unique ID
- Ideally - distribute towns to controllers evenly?

Should we even do this?

# **Learning Objectives for this Lesson**

**By the end of this lesson, you should be able to...**

- Characterize the benefits of replication and partitioning in distributed systems
- Evaluate the tradeoffs between consistency and availability in distributed systems

# This work is licensed under a Creative Commons Attribution-ShareAlike license

- This work is licensed under the Creative Commons Attribution-ShareAlike 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/4.0/>
- You are free to:
  - Share — copy and redistribute the material in any medium or format
  - Adapt — remix, transform, and build upon the material
    - for any purpose, even commercially.
- Under the following terms:
  - Attribution — You must give appropriate credit, provide a link to the license, and indicate if changes were made. You may do so in any reasonable manner, but not in any way that suggests the licensor endorses you or your use.
  - ShareAlike — If you remix, transform, or build upon the material, you must distribute your contributions under the same license as the original.
  - No additional restrictions — You may not apply legal terms or technological measures that legally restrict others from doing anything the license permits.